


# CS 325 I - Computer Networking I: Web and FTP

Professor Patrick Traynor  
Lecture 04  
8/29/2013


# Reminders

- The first project has been posted on the website.
  - Check the calendar for links! Due 9/12!
- Enter strings into the client and have the server return a special value for each one.
- Your job is to take care of the networking portion. I have already written the function to create the “special” value.
- Homework 1 is due at the beginning of the next class!




# Review

- Last time, we talked about principles of network applications
  - End-to-end argument
  - Network architectures (Client/Server, P2P)
  - Service requirements




# More Info: Bandwidth-Delay Product

- A number of students emailed questions about the “bandwidth-delay product”.
- This is simply the bandwidth of a link multiplied by the end-to-end delay (in seconds).
  - It tells us how many bits are “in flight”.
- Example: If we have a 10Mbps link between here and Berkeley (with a 100ms delay), what is the bandwidth-delay product?
  - $10\text{Mbps} * 1/10\text{sec} = 1 \text{ Mb}$



# Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- 2.3 FTP
- 2.4 Electronic Mail
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7-2.8 Sockets



# Web and HTTP

## First, some vocabulary:

- **Web page** consists of **objects**
- Object can be HTML file, JPEG image, Java applet, audio file, ...
- Web page consists of **base HTML-file** which includes several referenced objects
- Each object is addressable by a **URL**
- **Example URL:**  

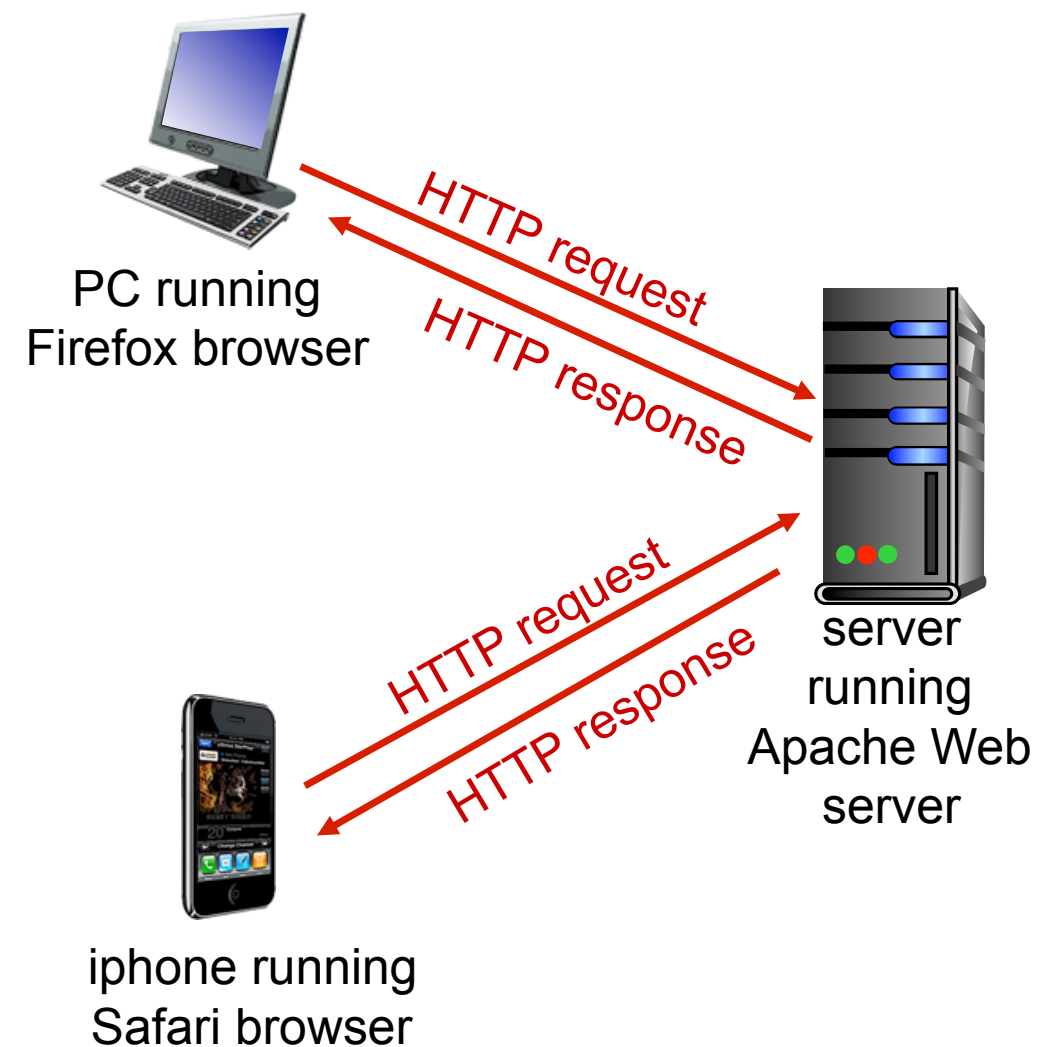
www.someschool.edu  
host name

/someDept/pic.gif  
path name

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - **client**: browser that requests, receives, “displays” Web objects
  - **server**: Web server sends objects in response to requests





# HTTP overview (continued)

## Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is “stateless”

- server maintains no information about past client requests

aside

## Protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled



# HTTP connections

## Nonpersistent HTTP

- At most one object is sent over a TCP connection.
- HTTP/1.0 uses nonpersistent HTTP

## Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server.
- HTTP/1.1 uses persistent connections in default mode



# Nonpersistent HTTP

Suppose user enters URL

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

**1a.** HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

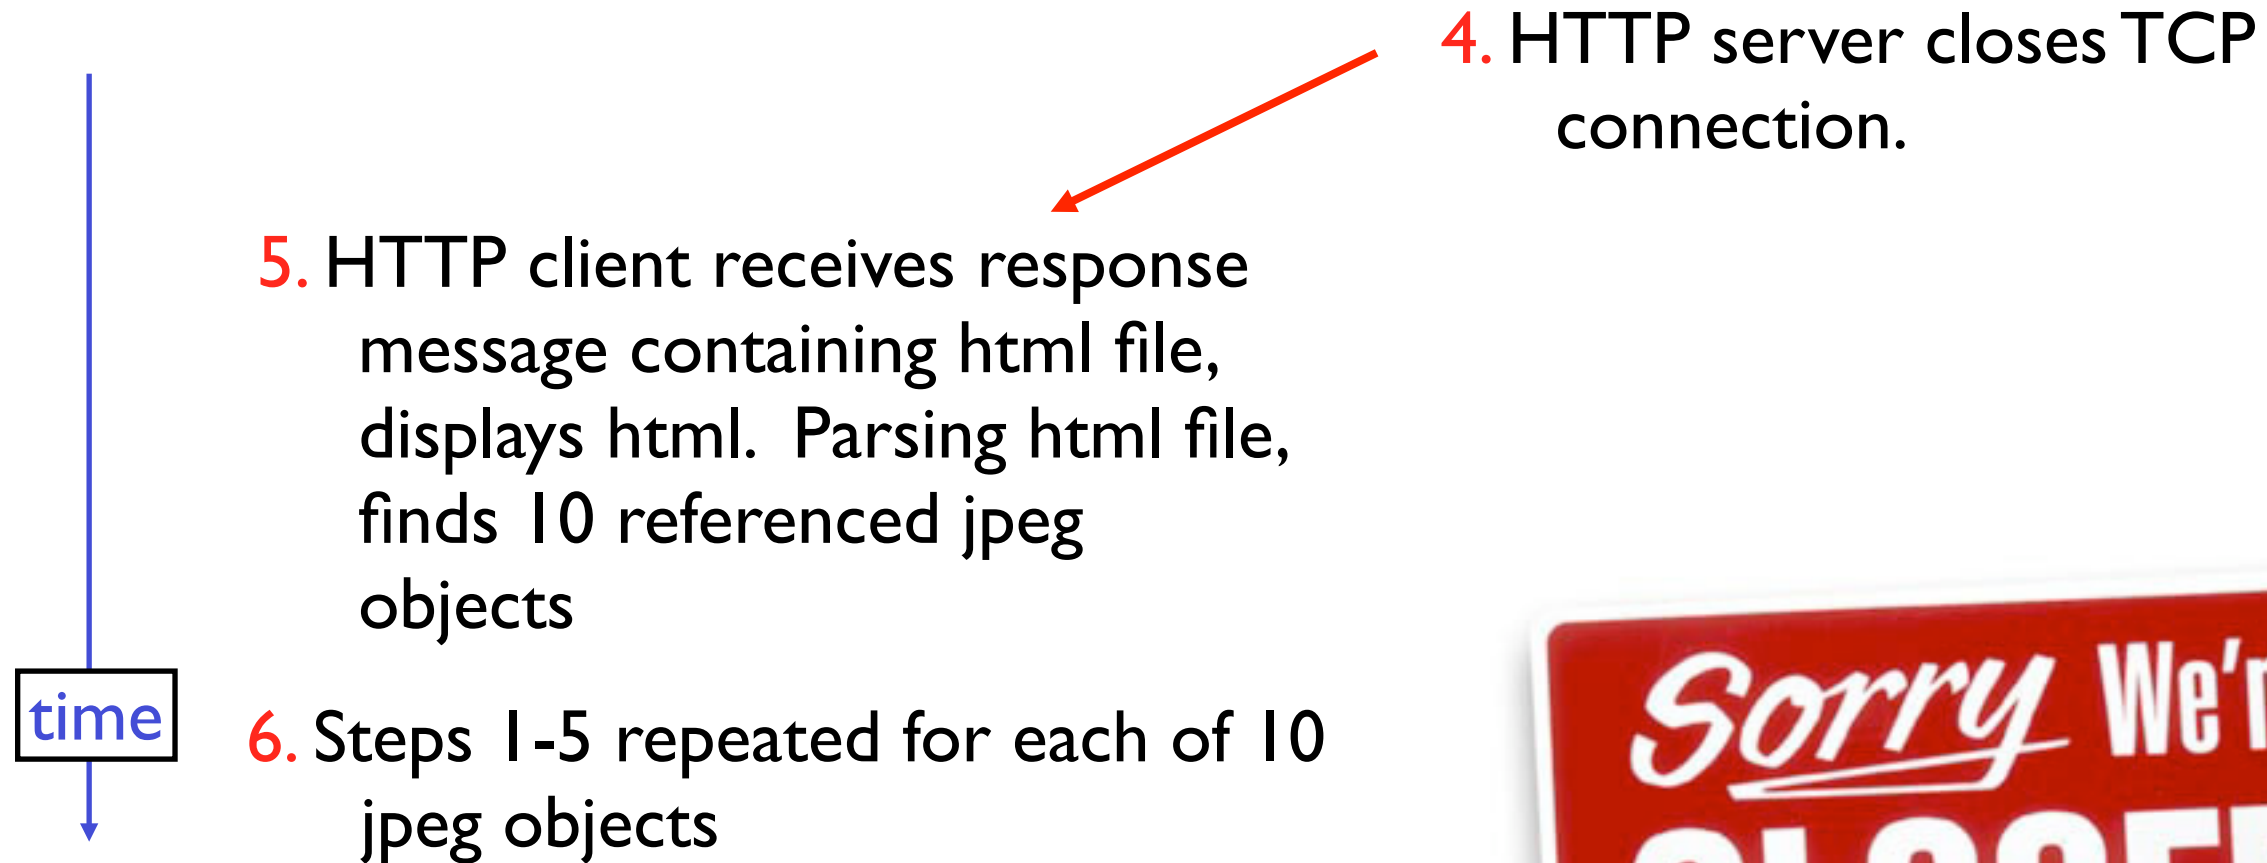
**1b.** HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

**2.** HTTP client sends HTTP **request message** (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

**3.** HTTP server receives request message, forms **response message** containing requested object, and sends message into its socket

time

# Nonpersistent HTTP (cont.)



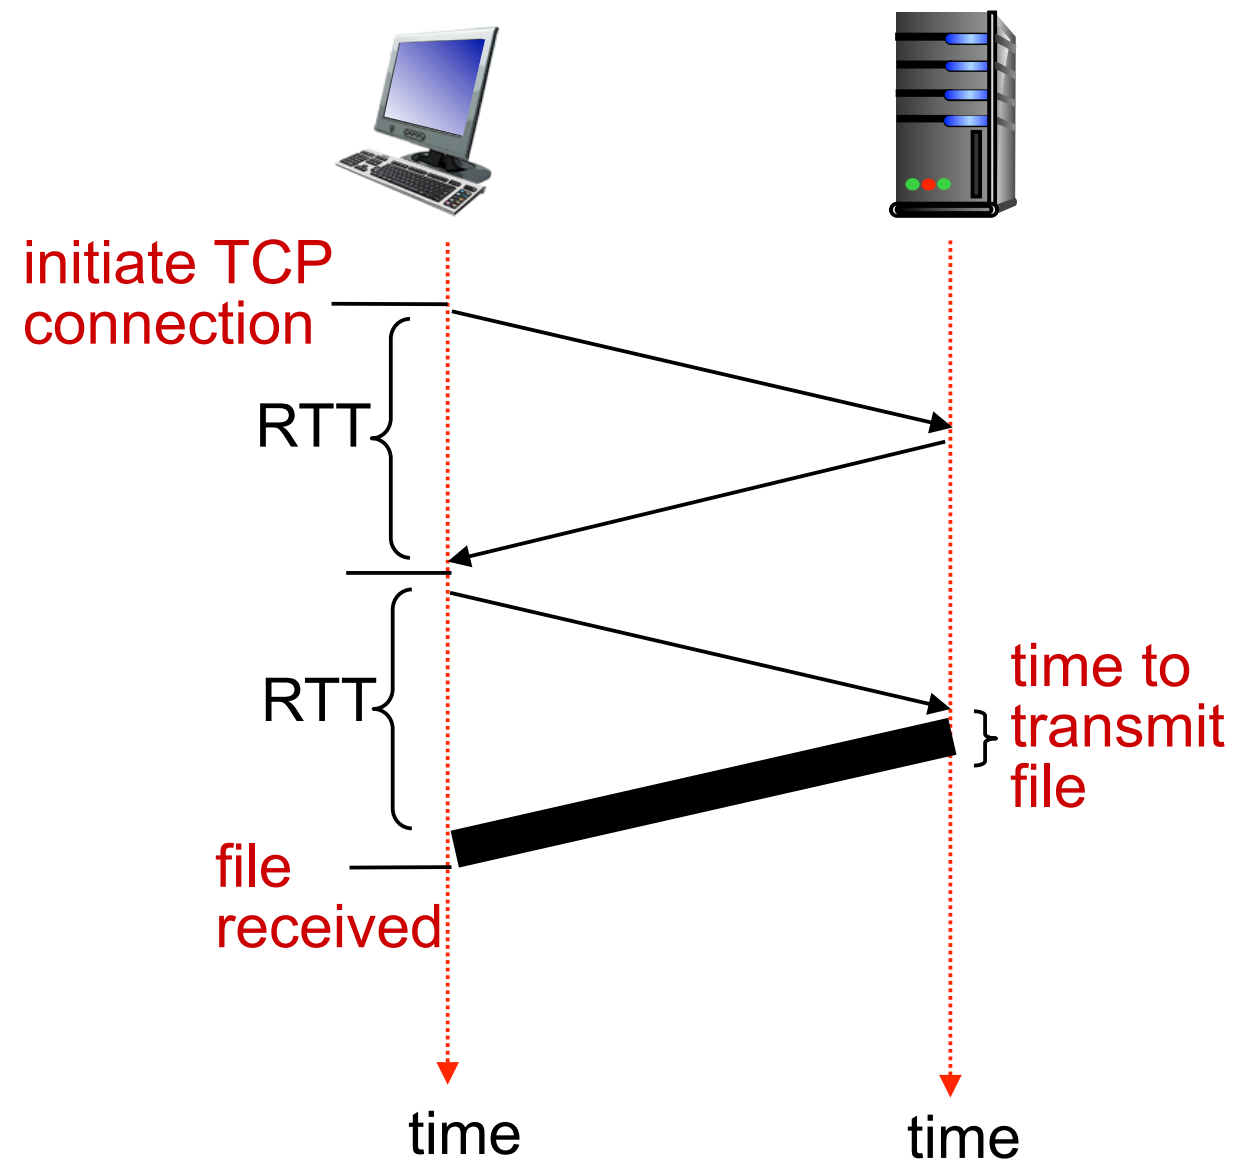
# Non-Persistent HTTP: Response time

**Definition of RTT:** time to send a small packet to travel from client to server and back.

## Response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

**total =  $2RTT + \text{transmit time}$**



# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

## Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

## Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

# Method types

## HTTP/1.0

- GET
- POST
- HEAD
  - asks server to leave requested object out of response


## HTTP/1.1

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# Uploading form input

## Post method:

- Web page often includes form input
- Input is uploaded to server in entity body



## URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`



# HTTP request message

- two types of HTTP messages: **request, response**
- **HTTP request message:**
  - ASCII (human-readable format)

The diagram illustrates the structure of an HTTP request message. It shows a request line followed by several header lines, each ending with a carriage return and line feed character. The annotations identify the components and control characters:

- request line (GET, POST, HEAD commands)**: Points to the first line of the message.
- header lines**: Points to the lines following the request line.
- carriage return, line feed at start of line indicates end of header lines**: Points to the final line of the header section.
- carriage return character**: Points to the `\r` character in the request line.
- line-feed character**: Points to the `\n` character in the request line.

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

# HTTP response message

status line  
(protocol  
status code  
status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT
\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html;
    charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

# HTTP response status codes

In first line in server to client response message.

A few sample codes:

## 200 OK

- request succeeded, requested object later in this message

## 301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

## 400 Bad Request

- request message not understood by server

## 404 Not Found

- requested document not found on this server

## 505 HTTP Version Not Supported

# Trying out HTTP (client side) for yourself

## 1. Telnet to your favorite Web server:

```
telnet www.cc.gatech.edu 80
```

Opens TCP connection to port 80  
(default HTTP server port) at  
`www.cc.gatech.edu`.  
Anything typed in sent  
to port 80 at www.cc.gatech.edu

## 2. Type in a GET HTTP request:

```
GET /~traynor/cs3251/f13/ HTTP/1.1  
Host: www.cc.gatech.edu
```

By typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

## 3. Look at response message sent by HTTP server!

# User-server state: cookies

Many major Web sites use cookies

## Four components:

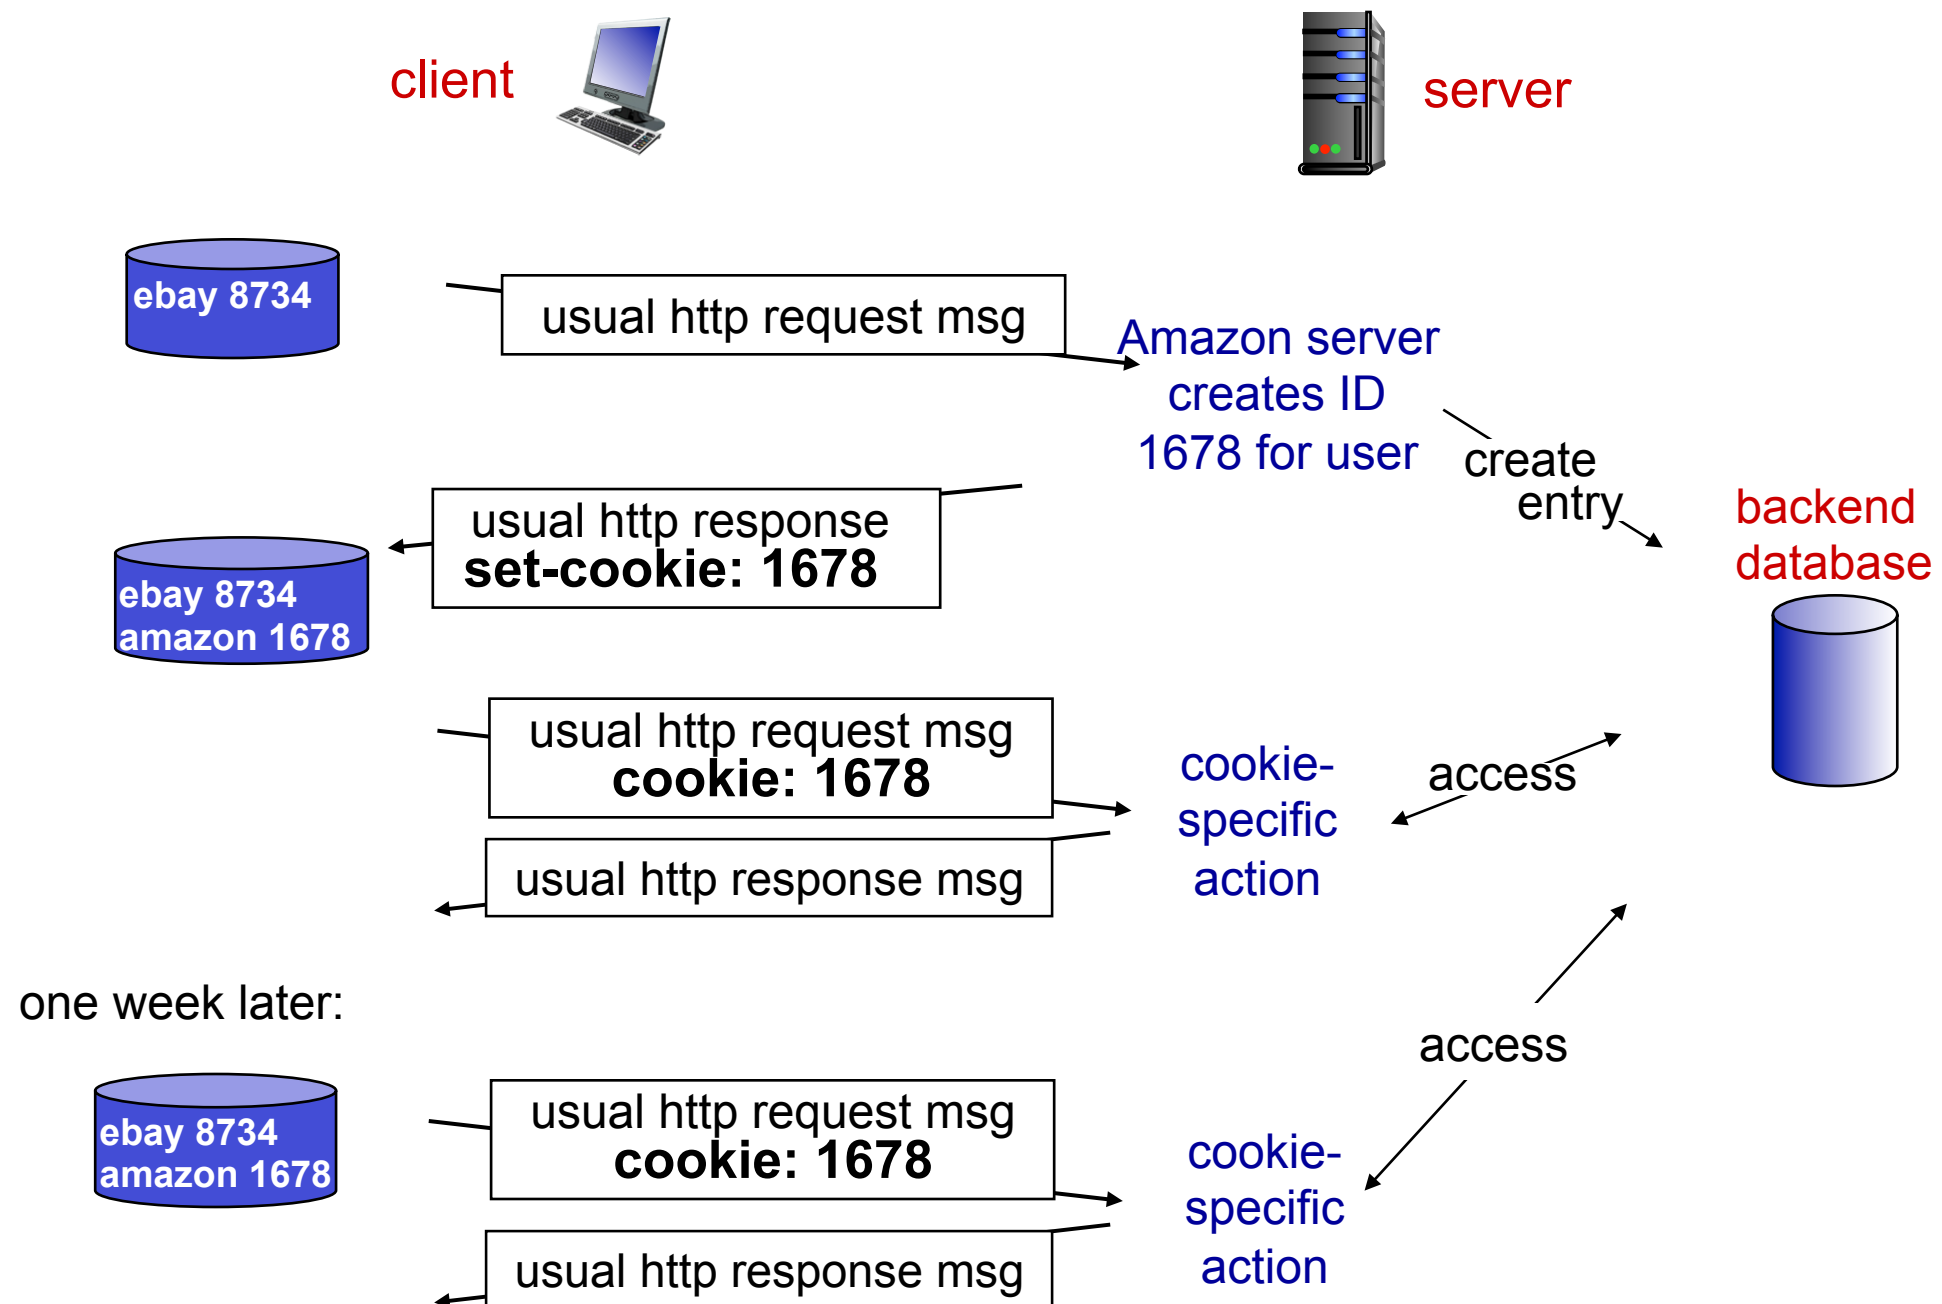
- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

## Example:

- ▶ Susan access Internet always from same PC
- ▶ She visits a specific e-commerce site for first time
- ▶ When initial HTTP requests arrives at site, site creates a unique ID and creates an entry in backend database for ID



# Cookies: keeping “state” (cont.)



# Cookies (continued)

## What cookies can bring:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## How to keep “state”:

- Protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

aside

## Cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites






# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client



# More about Web caching

- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)



## Why Web caching?

- Reduce response time for client request.
- Reduce traffic on an institution's access link.
- Internet dense with caches: enables “poor” content providers to effectively deliver content (but so does P2P file sharing)


# Caching example

## Assumptions

- average object size = 100,000 bits (100 Kb)
- avg. request rate from institution's browsers to origin servers = 15/sec
- avg data rate to browsers: 1.50 Mbps
- access link rate: 1.54 Mbps
- RTT from institutional router to any origin server: 2 sec

## Consequences

- utilization on LAN = 0.15%
- utilization on access link = **~97%**
- total delay = Internet delay + access delay + LAN delay  
  
= 2 sec + minutes + milliseconds




# Caching example (cont)

## Possible solution

- increase bandwidth of access link to, say, 154 Mbps

## Consequences

- utilization on LAN = 0.15%
- utilization on access link = ~1%
- Total delay = Internet delay + access delay + LAN delay  
= 2 sec + msec + msec
- So... how much is this going to cost us?




# Caching example (cont)

## Install cache

- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin

## Consequence

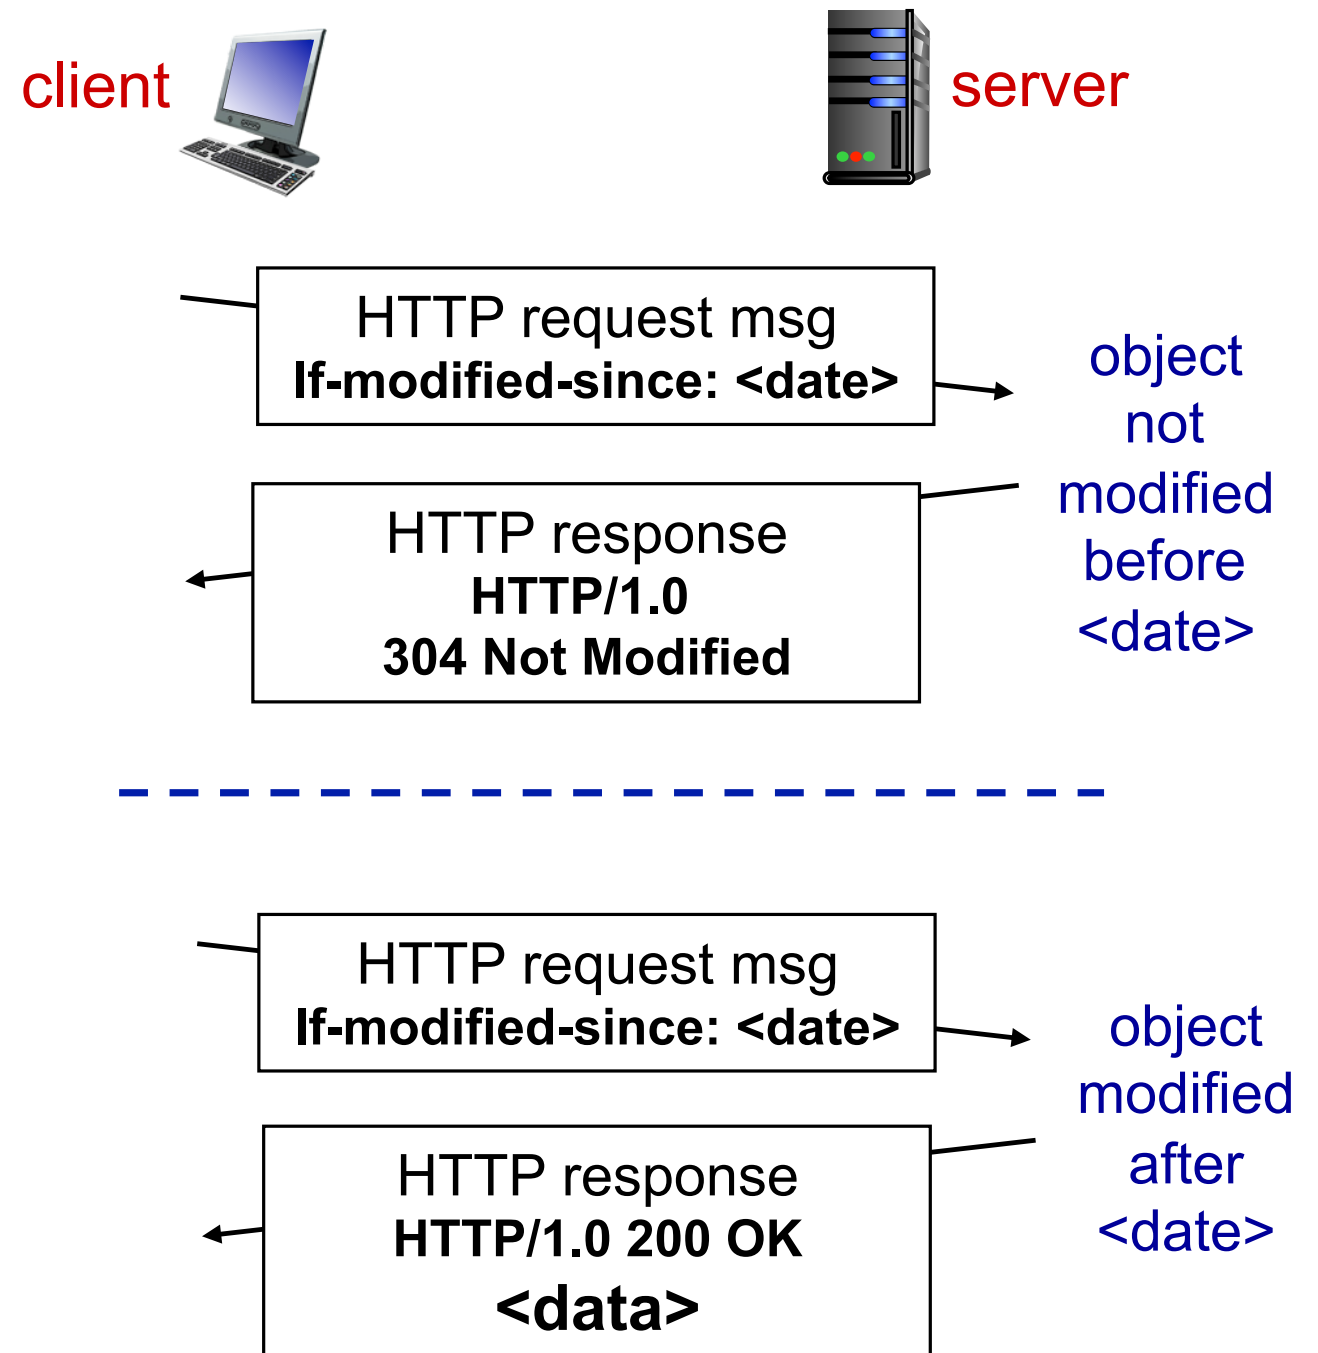
- access link utilization:
  - 60% of requests use access link
- data rate to browsers over access link =  $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - utilization =  $0.9 / 1.54 = .58$
- total delay
  - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
- $= 0.6 (2.01) + 0.4 (\sim \text{msecs})$
- $= \sim 1.2 \text{ secs}$
- less than with 154 Mbps link (and cheaper too!)



# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request  
`If-modified-since: <date>`
- server: response contains no object if cached copy is up-to-date:

`HTTP/1.0 304 Not Modified`






# Chapter 2: Application layer

- 2.1 Principles of network applications
- 2.2 Web and HTTP
- **2.3 FTP**
- 2.4 Electronic Mail
- 2.5 DNS
- 2.6 P2P file sharing
- 2.7-2.8 Sockets





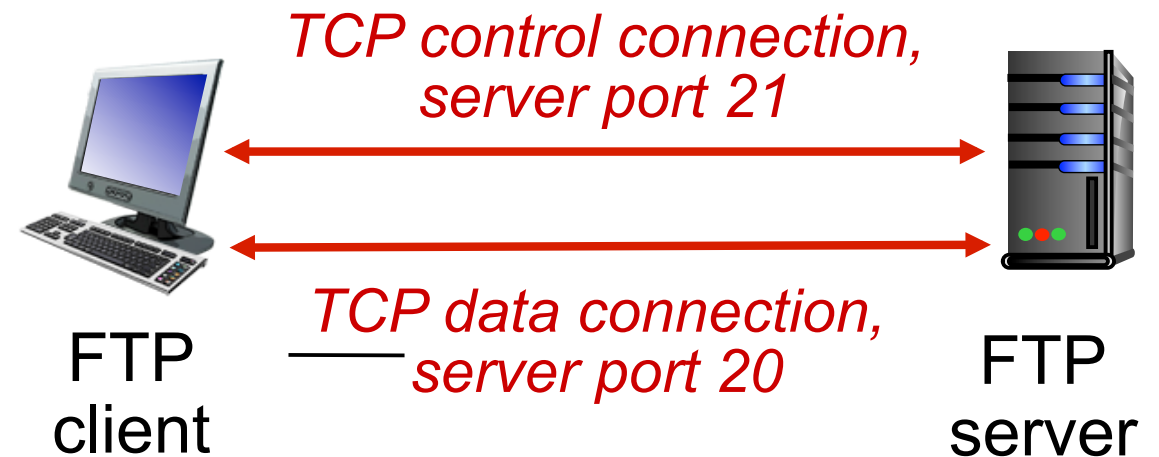
# FTP: the file transfer protocol



- transfer file to/from remote host
- client/server model
  - **client**: side that initiates transfer (either to/from remote)
  - **server**: remote host
- ftp: RFC 959
- ftp server: port 21

# FTP: separate control, data connections

- FTP client contacts FTP server at port 21, specifying TCP as transport protocol
- Client obtains authorization over control connection
- Client browses remote directory by sending commands over control connection.
- When server receives file transfer command, server opens 2<sup>nd</sup> TCP connection (for file) to client
- After transferring one file, server closes data connection.



- Server opens another TCP data connection to transfer another file.
- Control connection: “out of band”
- FTP server maintains “state”: current directory, earlier authentication

# FTP commands, responses

## Sample commands:

- sent as ASCII text over control channel
- **USER *username***
- **PASS *password***
- **LIST** return list of file in current directory
- **RETR *filename*** retrieves (gets) file
- **STOR *filename*** stores (puts) file onto remote host

## Sample return codes

- status code and phrase (as in HTTP)
- **331 Username OK, password required**
- **125 data connection already open; transfer starting**
- **425 Can't open data connection**
- **452 Error writing file**

# Next Time

- We will cover Email and DNS
  - Read Sections 2.4 and 2.5
- Reminder:
  - Project I has been posted
  - Homework I is due at the beginning of next class, turned in via T-Square **ONLY**.

Next is what?™  
SAMSUNG mobile