

# CS 325 I - Computer Networks I: Reliable Data Transfer

Professor Patrick Traynor  
Lecture 08  
9/12/13

# Announcements

- Project 1 is due today
  - Submit your tarball **via T-Square**.
  - Due by 5pm!
- Project 2 will be posted by Tuesday.
  - We will talk about the details in class on Tuesday.
  - Get started soon! This will require a significant amount of effort to get done.
- Homework 2 has been posted
  - Submission deadline pushed back to 10/1



# Last Time

- Multiplexing/Demultiplexing at the Transport Layer.
  - How do TCP and UDP differ?
- UDP gives us virtually “bare-bones” access to the network layer.
  - What are the four fields in a UDP header?
- What is port scanning?
  - How can it be used to protect systems? To attack them?



# Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
  - ▶ segment structure
  - ▶ reliable data transfer
  - ▶ flow control
  - ▶ connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

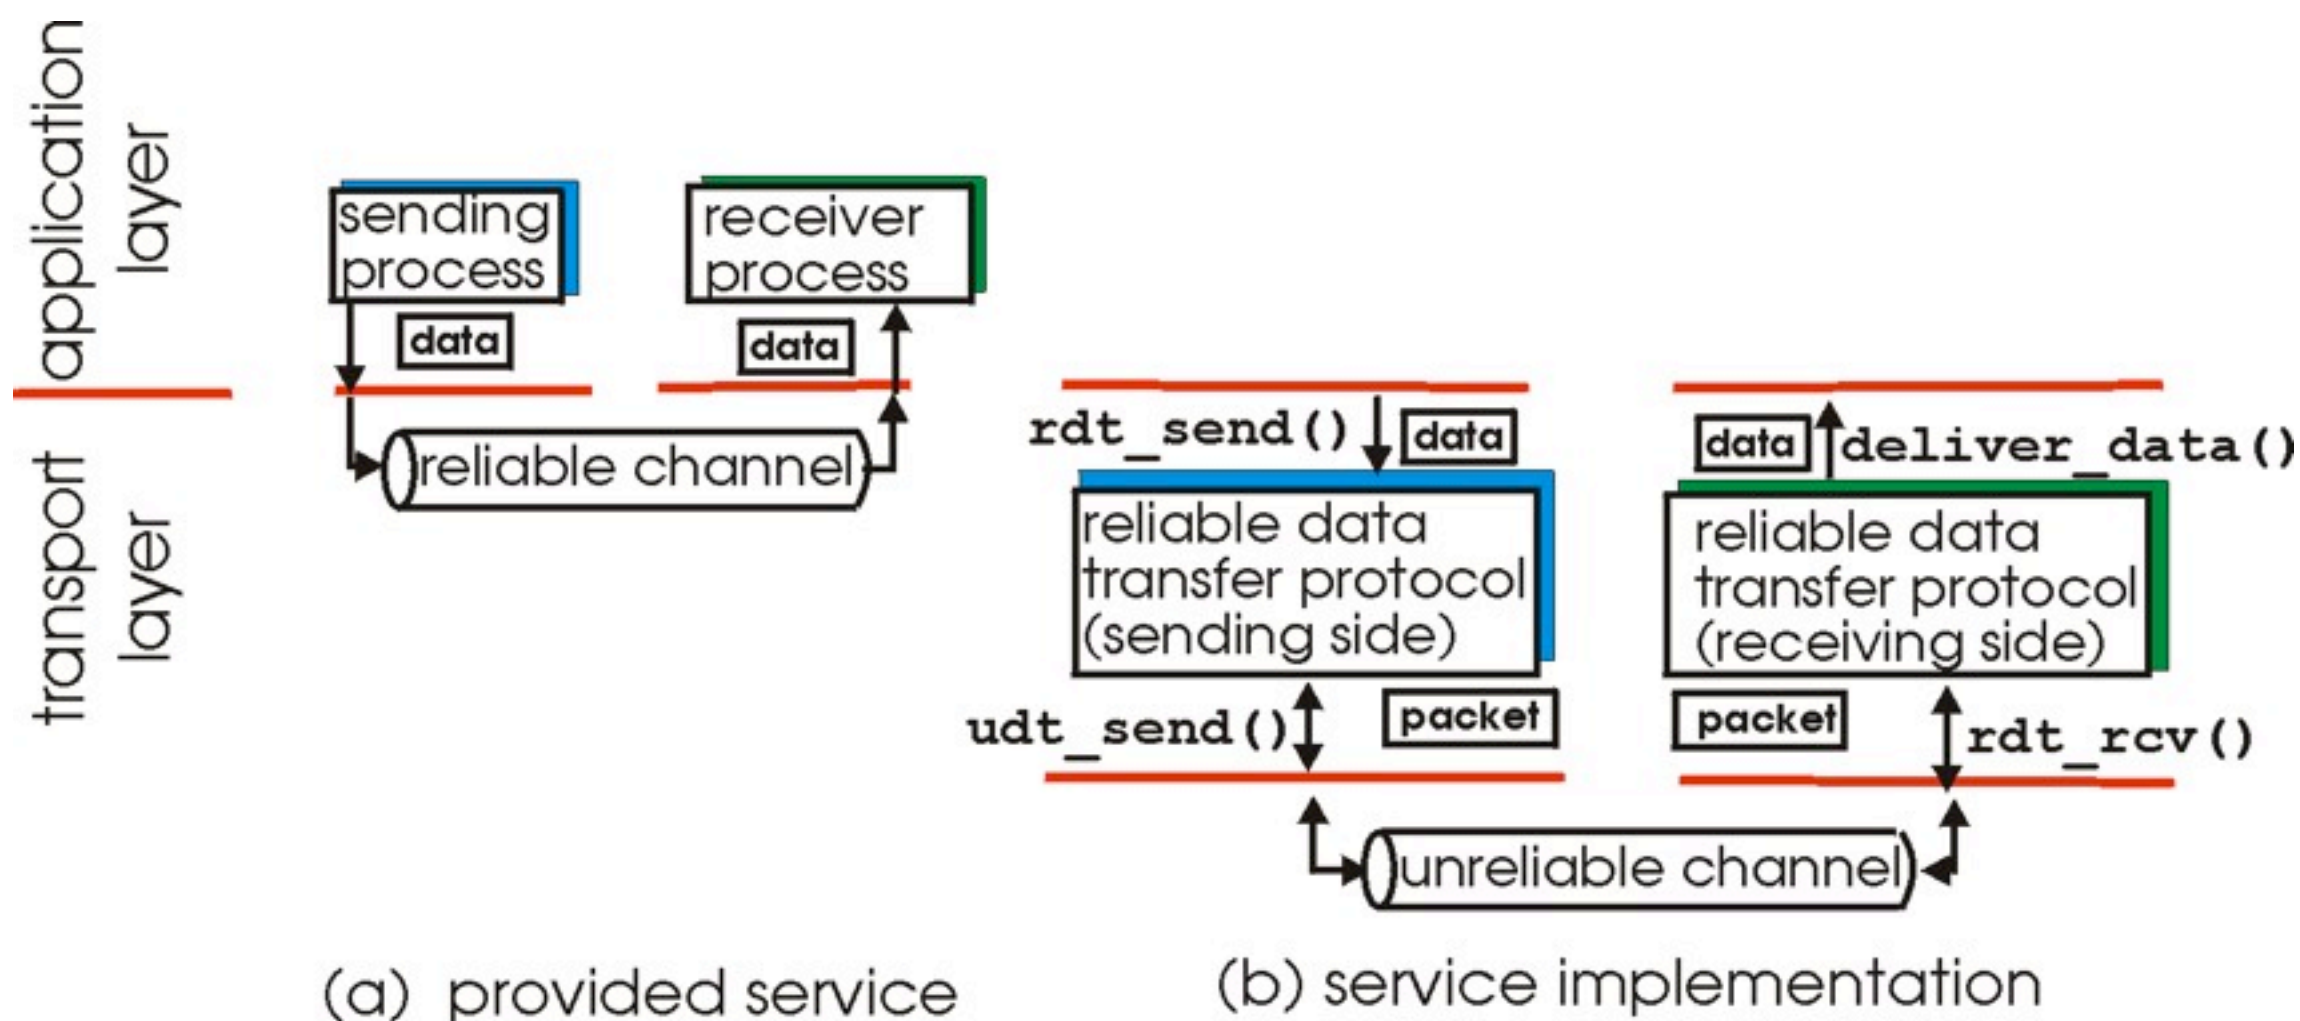
# Problem Solving

- Given all of our talk about TCP, you might assume that describing how it works is straightforward.
- The reality is that there are lots of different ways that we could have provided “reliable” delivery.
  - Here is where we will really start to see design tradeoffs.
- It’s time in your CS career to be creative...
  - Let’s solve some problems.



# Principles of Reliable data transfer

- important in app., transport, link layers
- top-10 list of important networking topics!

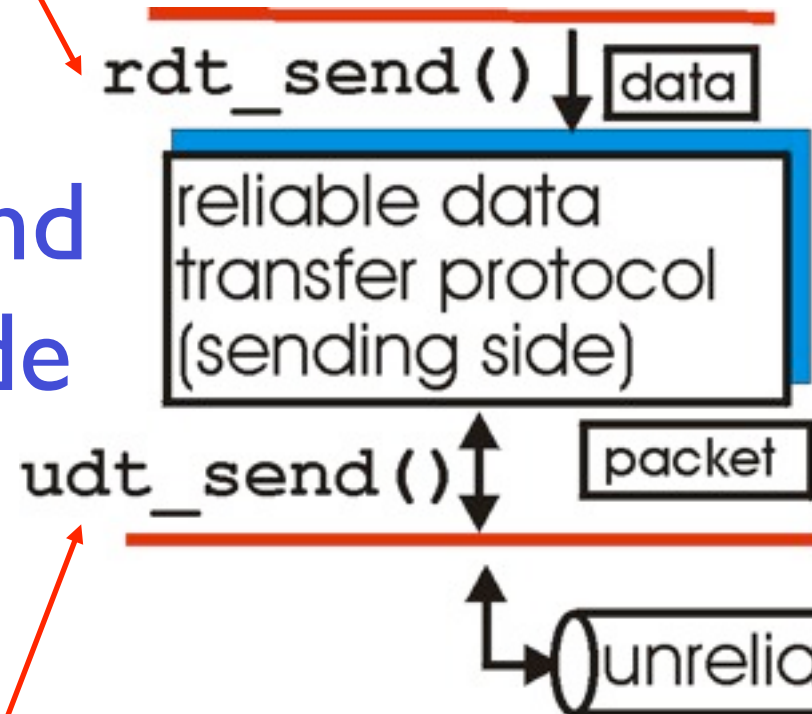


- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Reliable data transfer: getting started

**rdt\_send()** : called from above,  
(e.g., by app.). Passed data to  
deliver to receiver upper layer

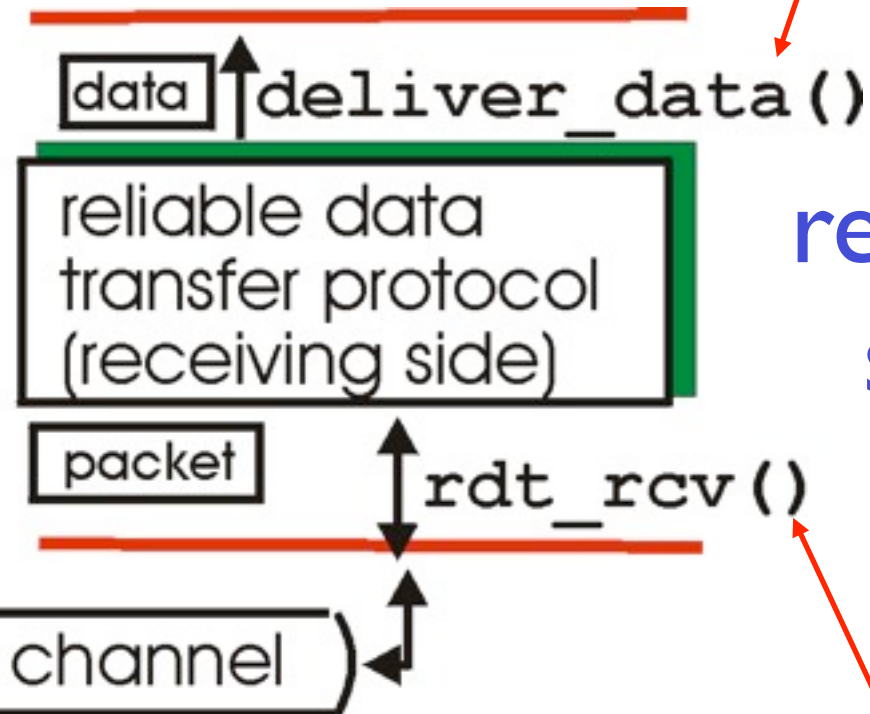
send  
side



**udt\_send()** : called by rdt,  
to transfer packet over  
unreliable channel to receiver

**deliver\_data()** : called by  
rdt to deliver data to upper

receive  
side



**rdt\_rcv()** : called when packet  
arrives on rcv-side of channel

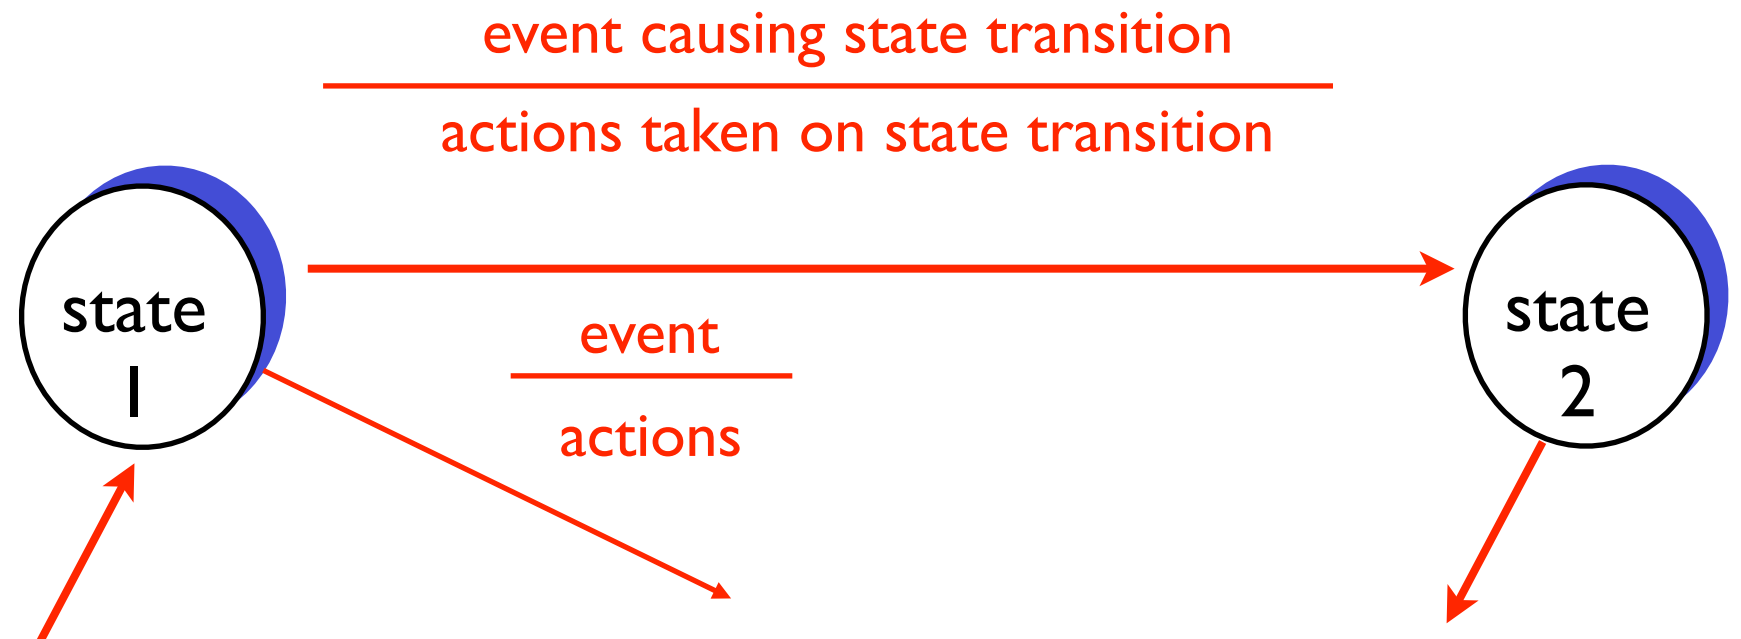


# Reliable data transfer: getting started

## We'll:

- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
  - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver

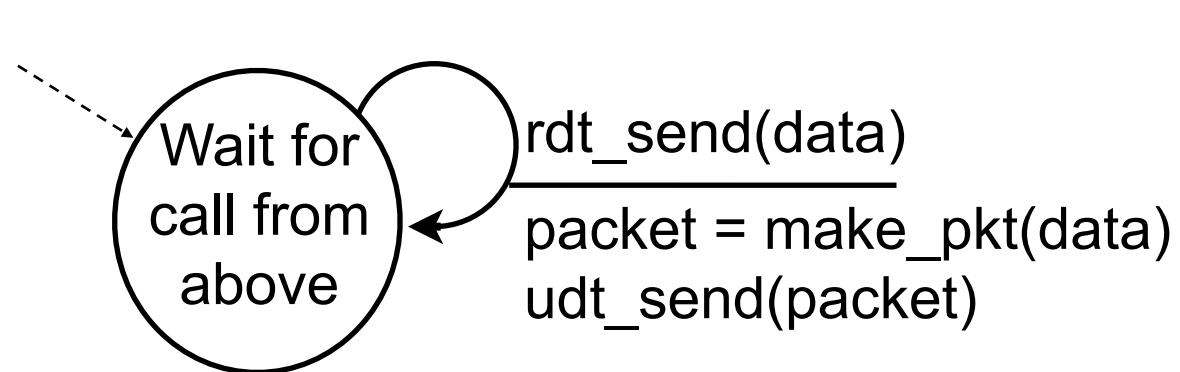
**state:** when in this “state”  
next state uniquely  
determined by next  
event






# Rdt 1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver read data from underlying channel



sender

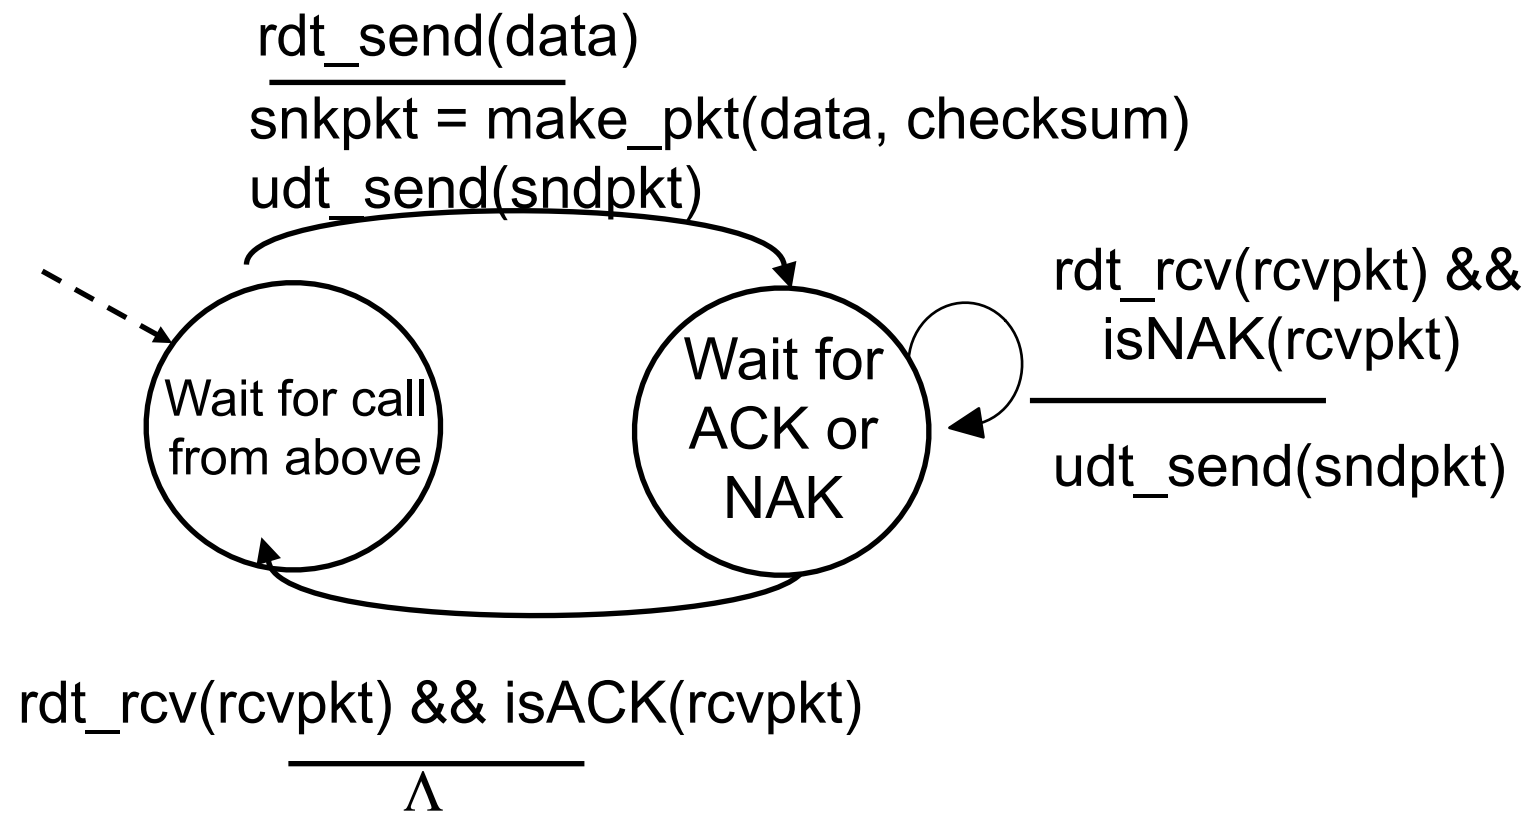


receiver

# Rdt2.0: channel with bit errors


- underlying channel may flip bits in packet
  - checksum to detect bit errors
- *the question: how to recover from errors:*
  - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
  - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK
- new mechanisms in **rdt2.0** (beyond **rdt1.0**):
  - error detection
  - receiver feedback: control msgs (ACK,NAK) rcvr->sender

# rdt2.0: FSM specification

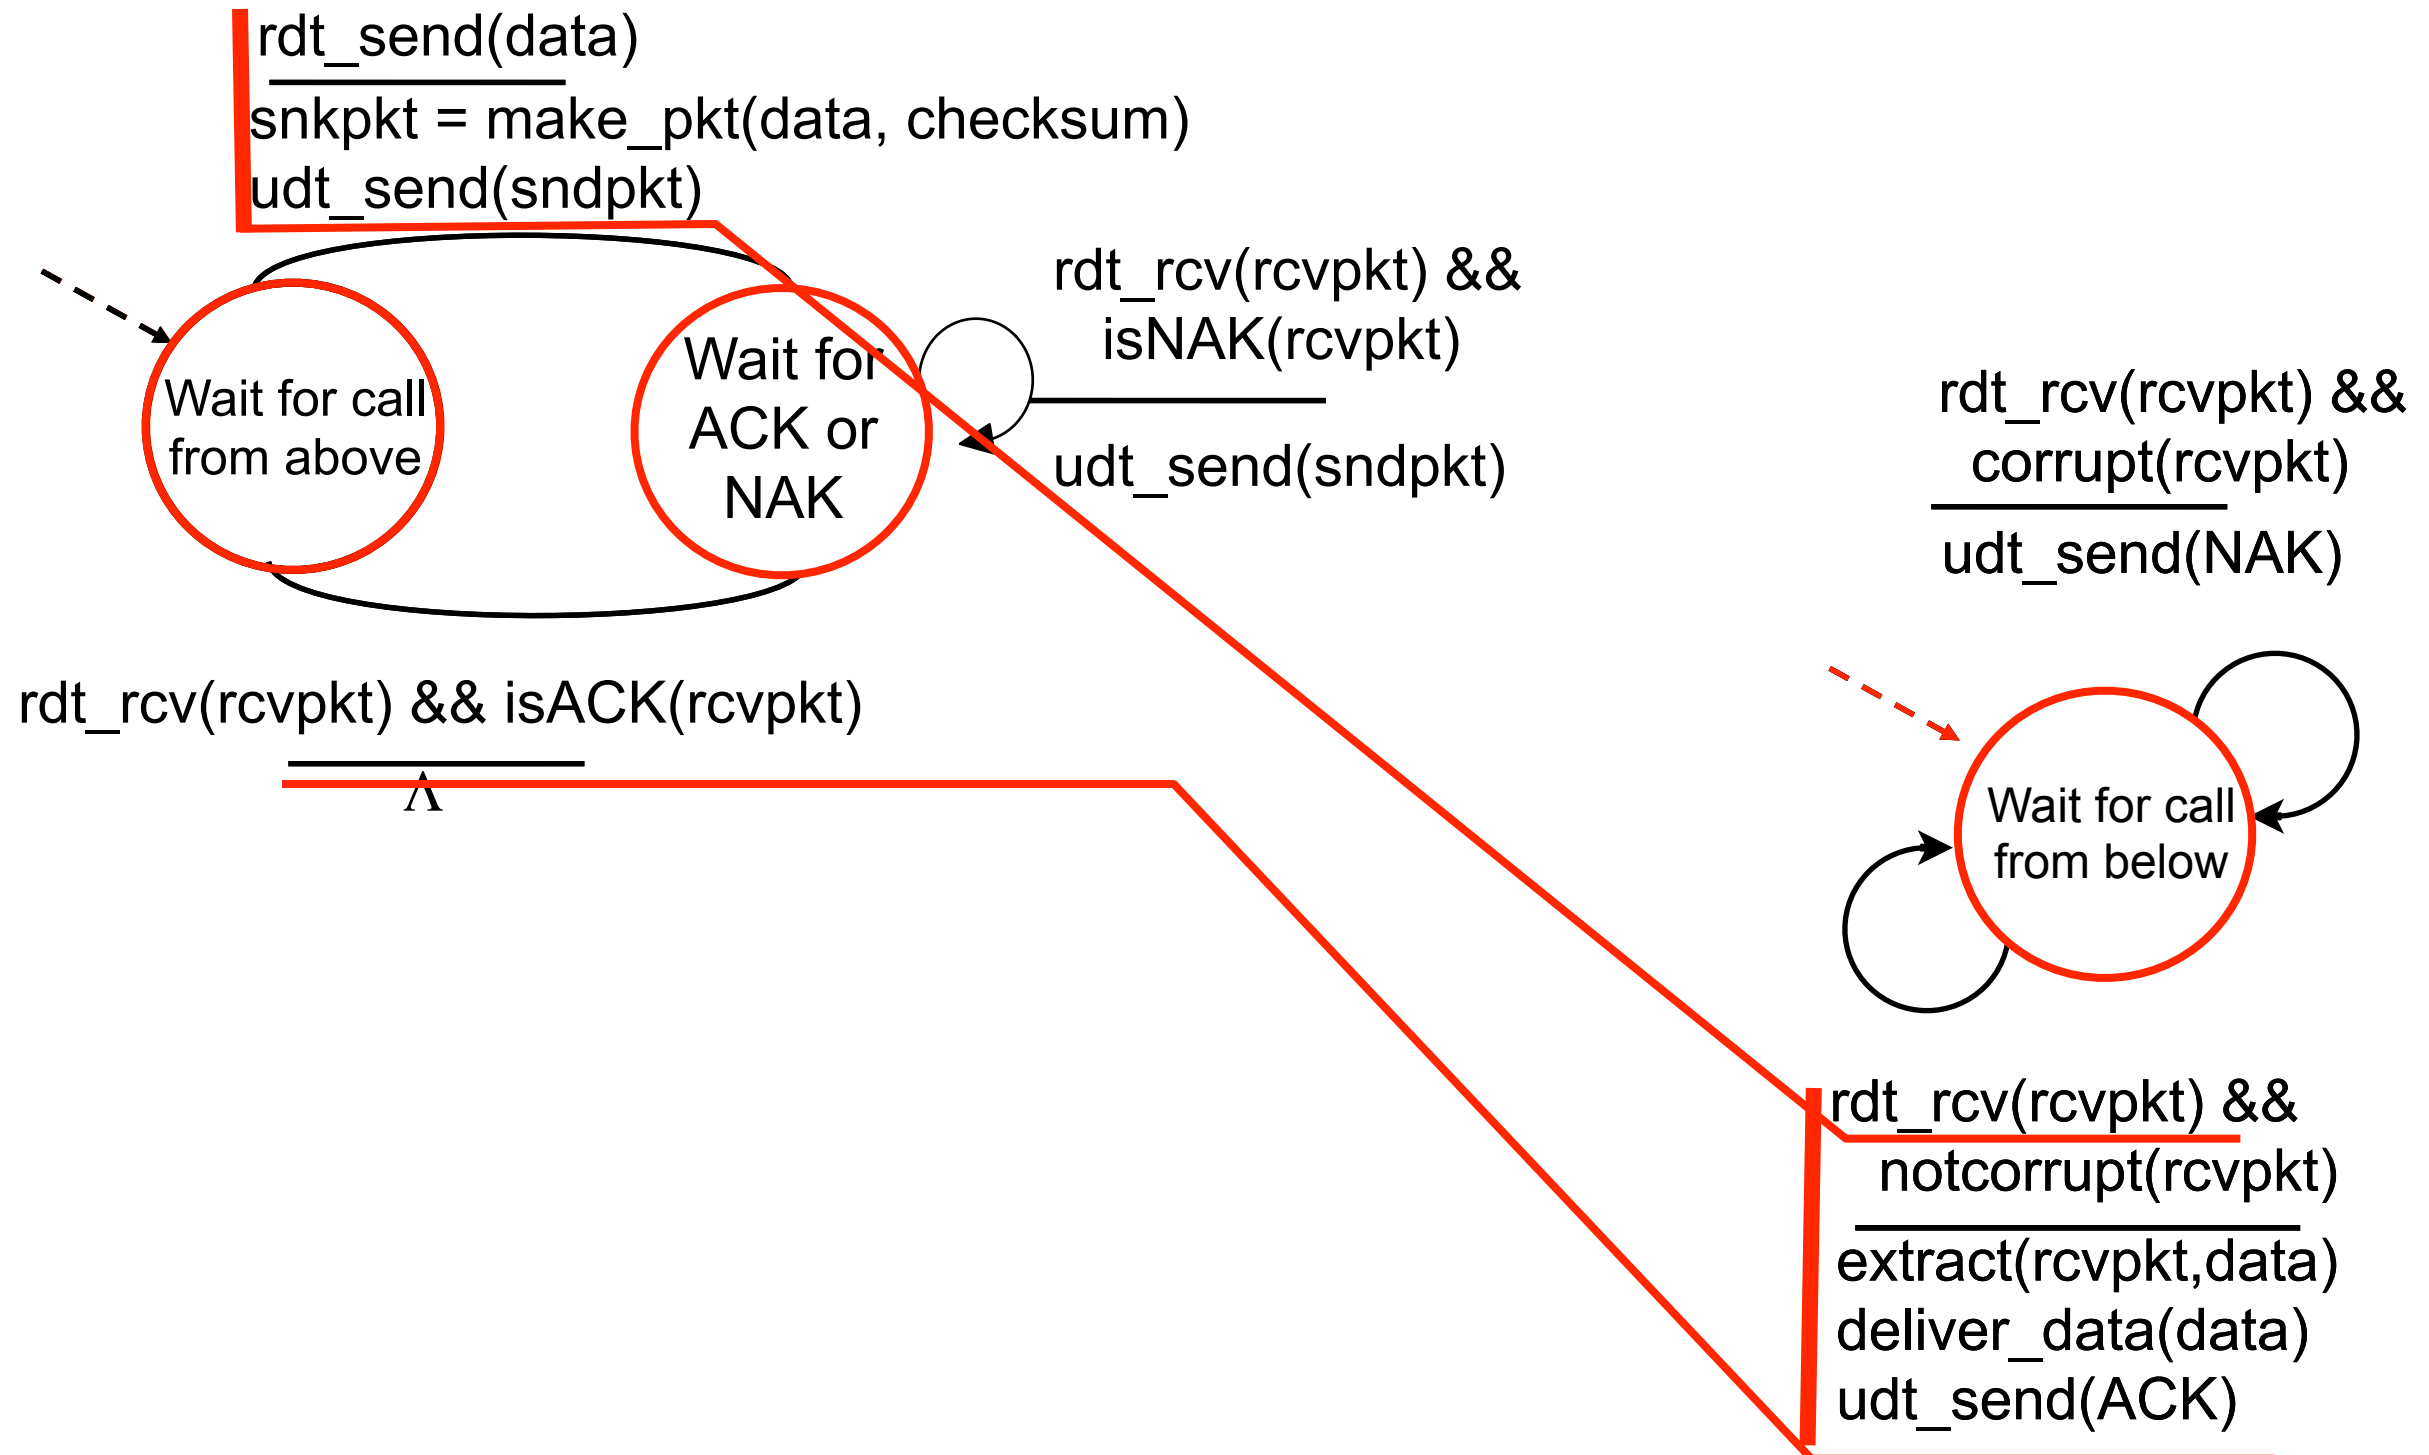


sender

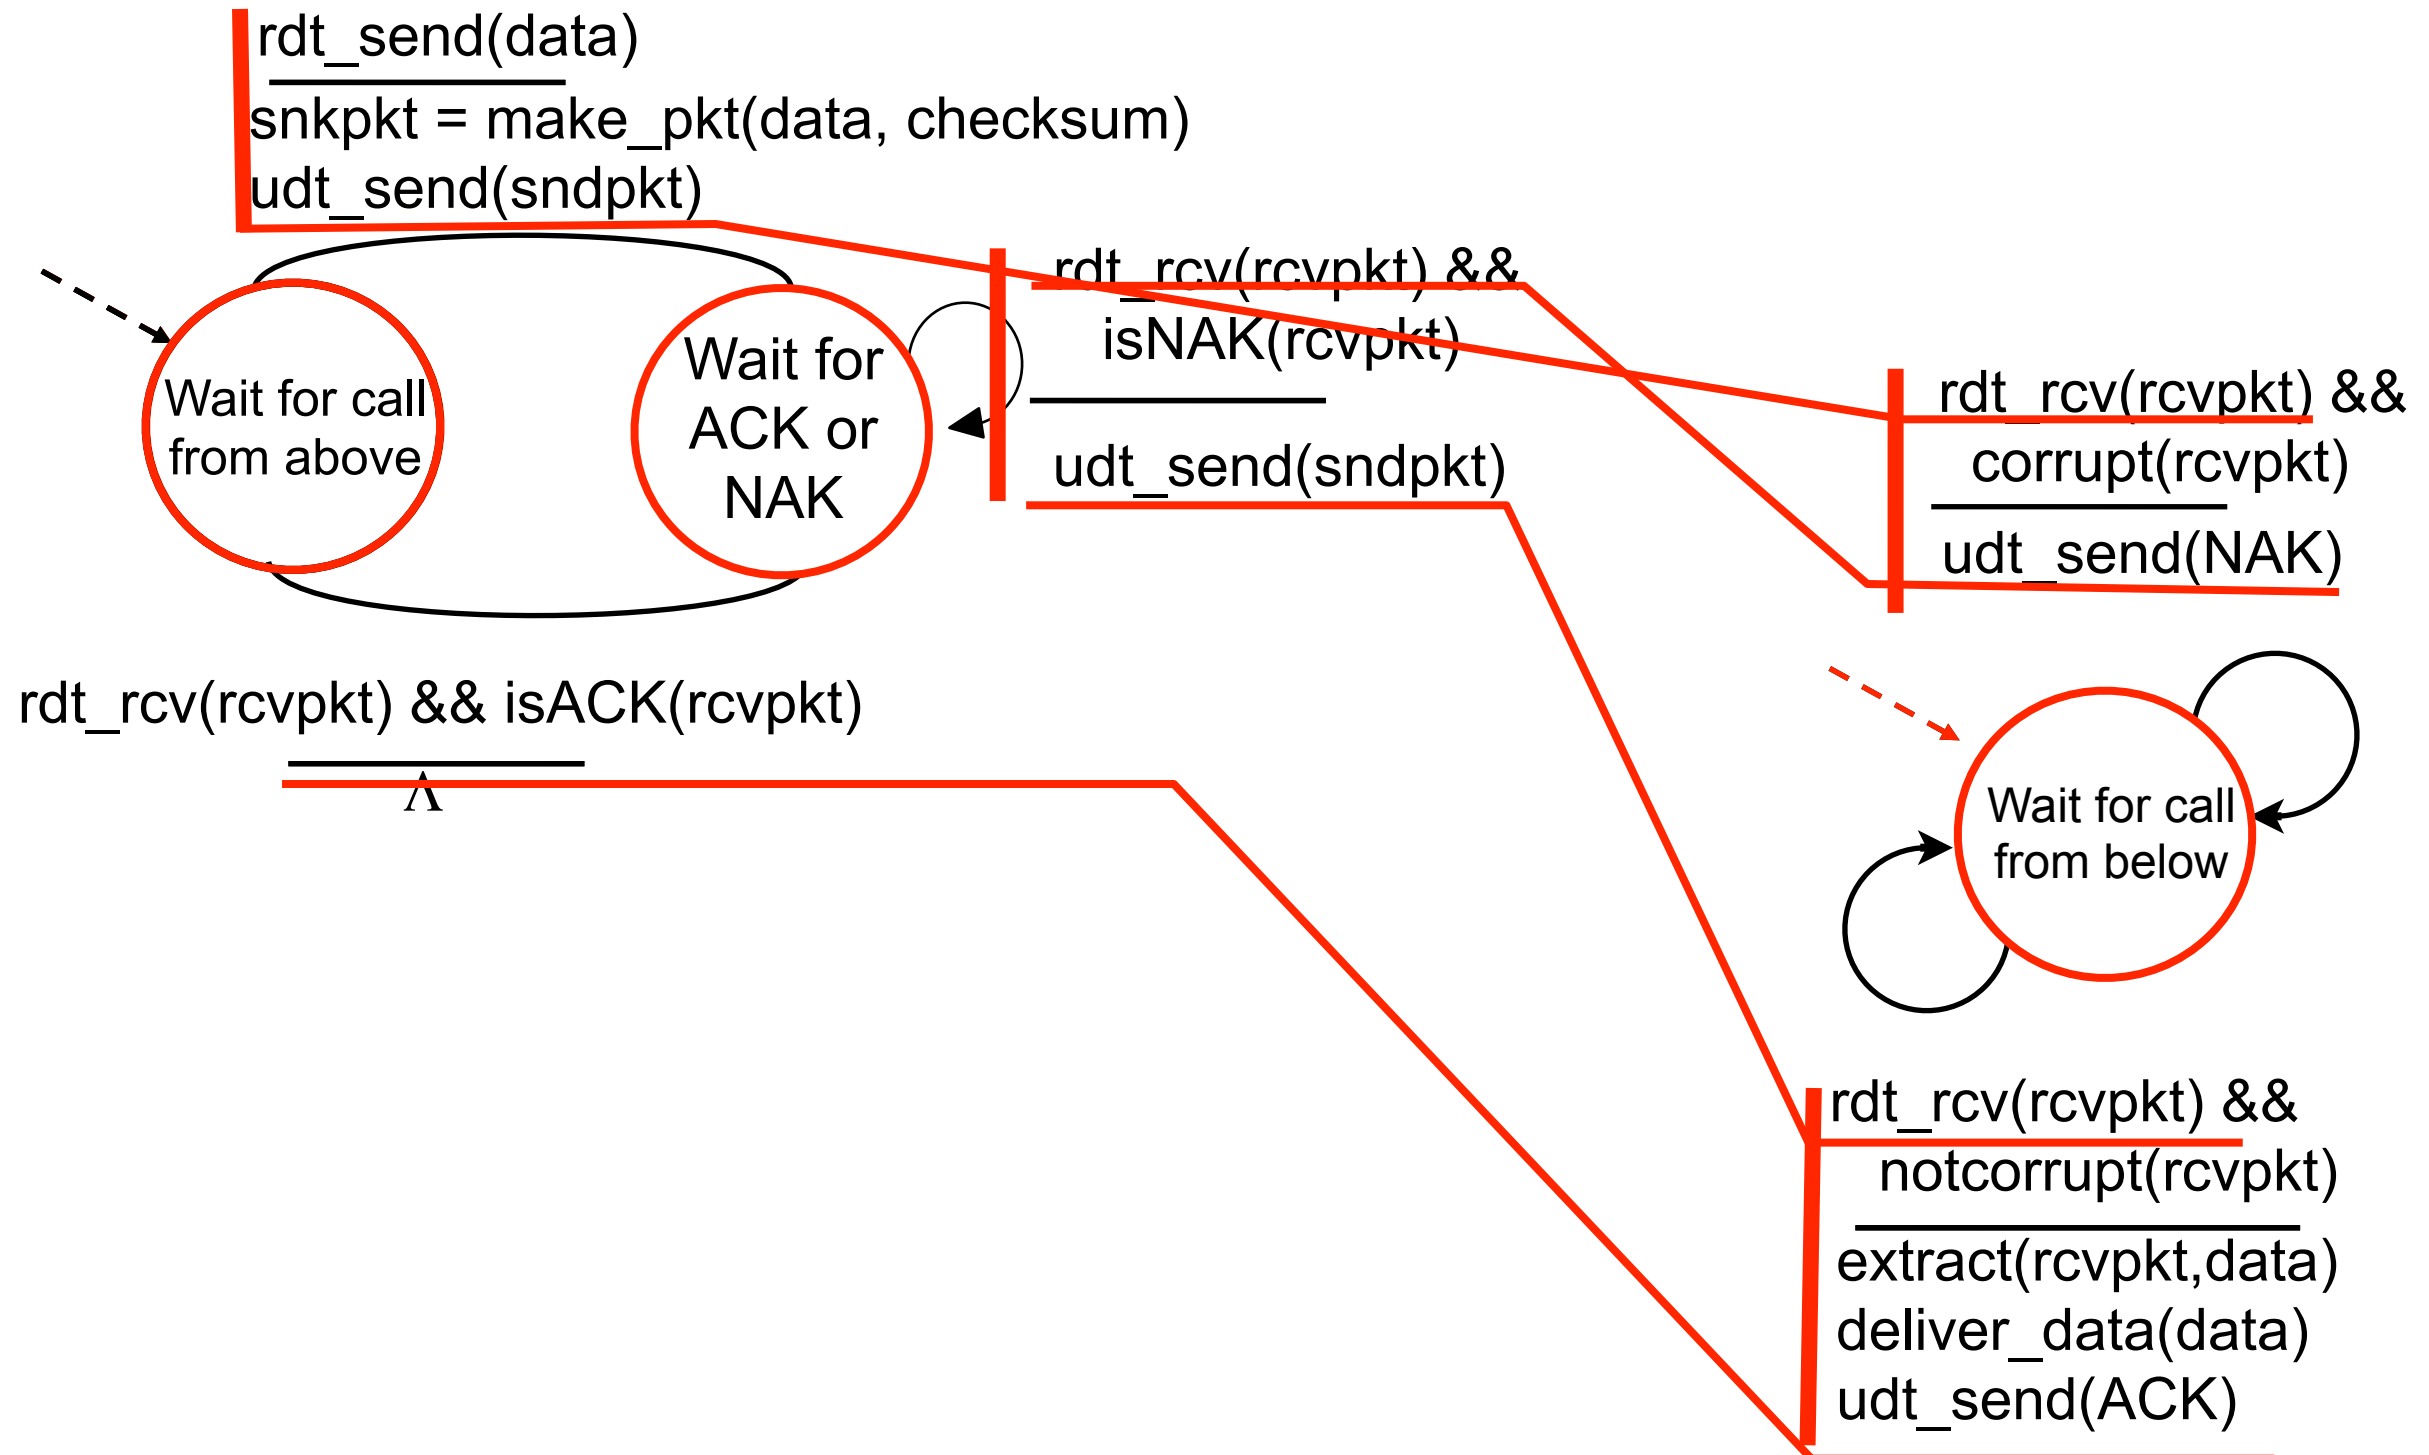
receiver



# rdt2.0: operation with no errors



# rdt2.0: error scenario




# rdt2.0 has a fatal flaw!

## What happens if ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit: possible duplicate

## Handling duplicates:

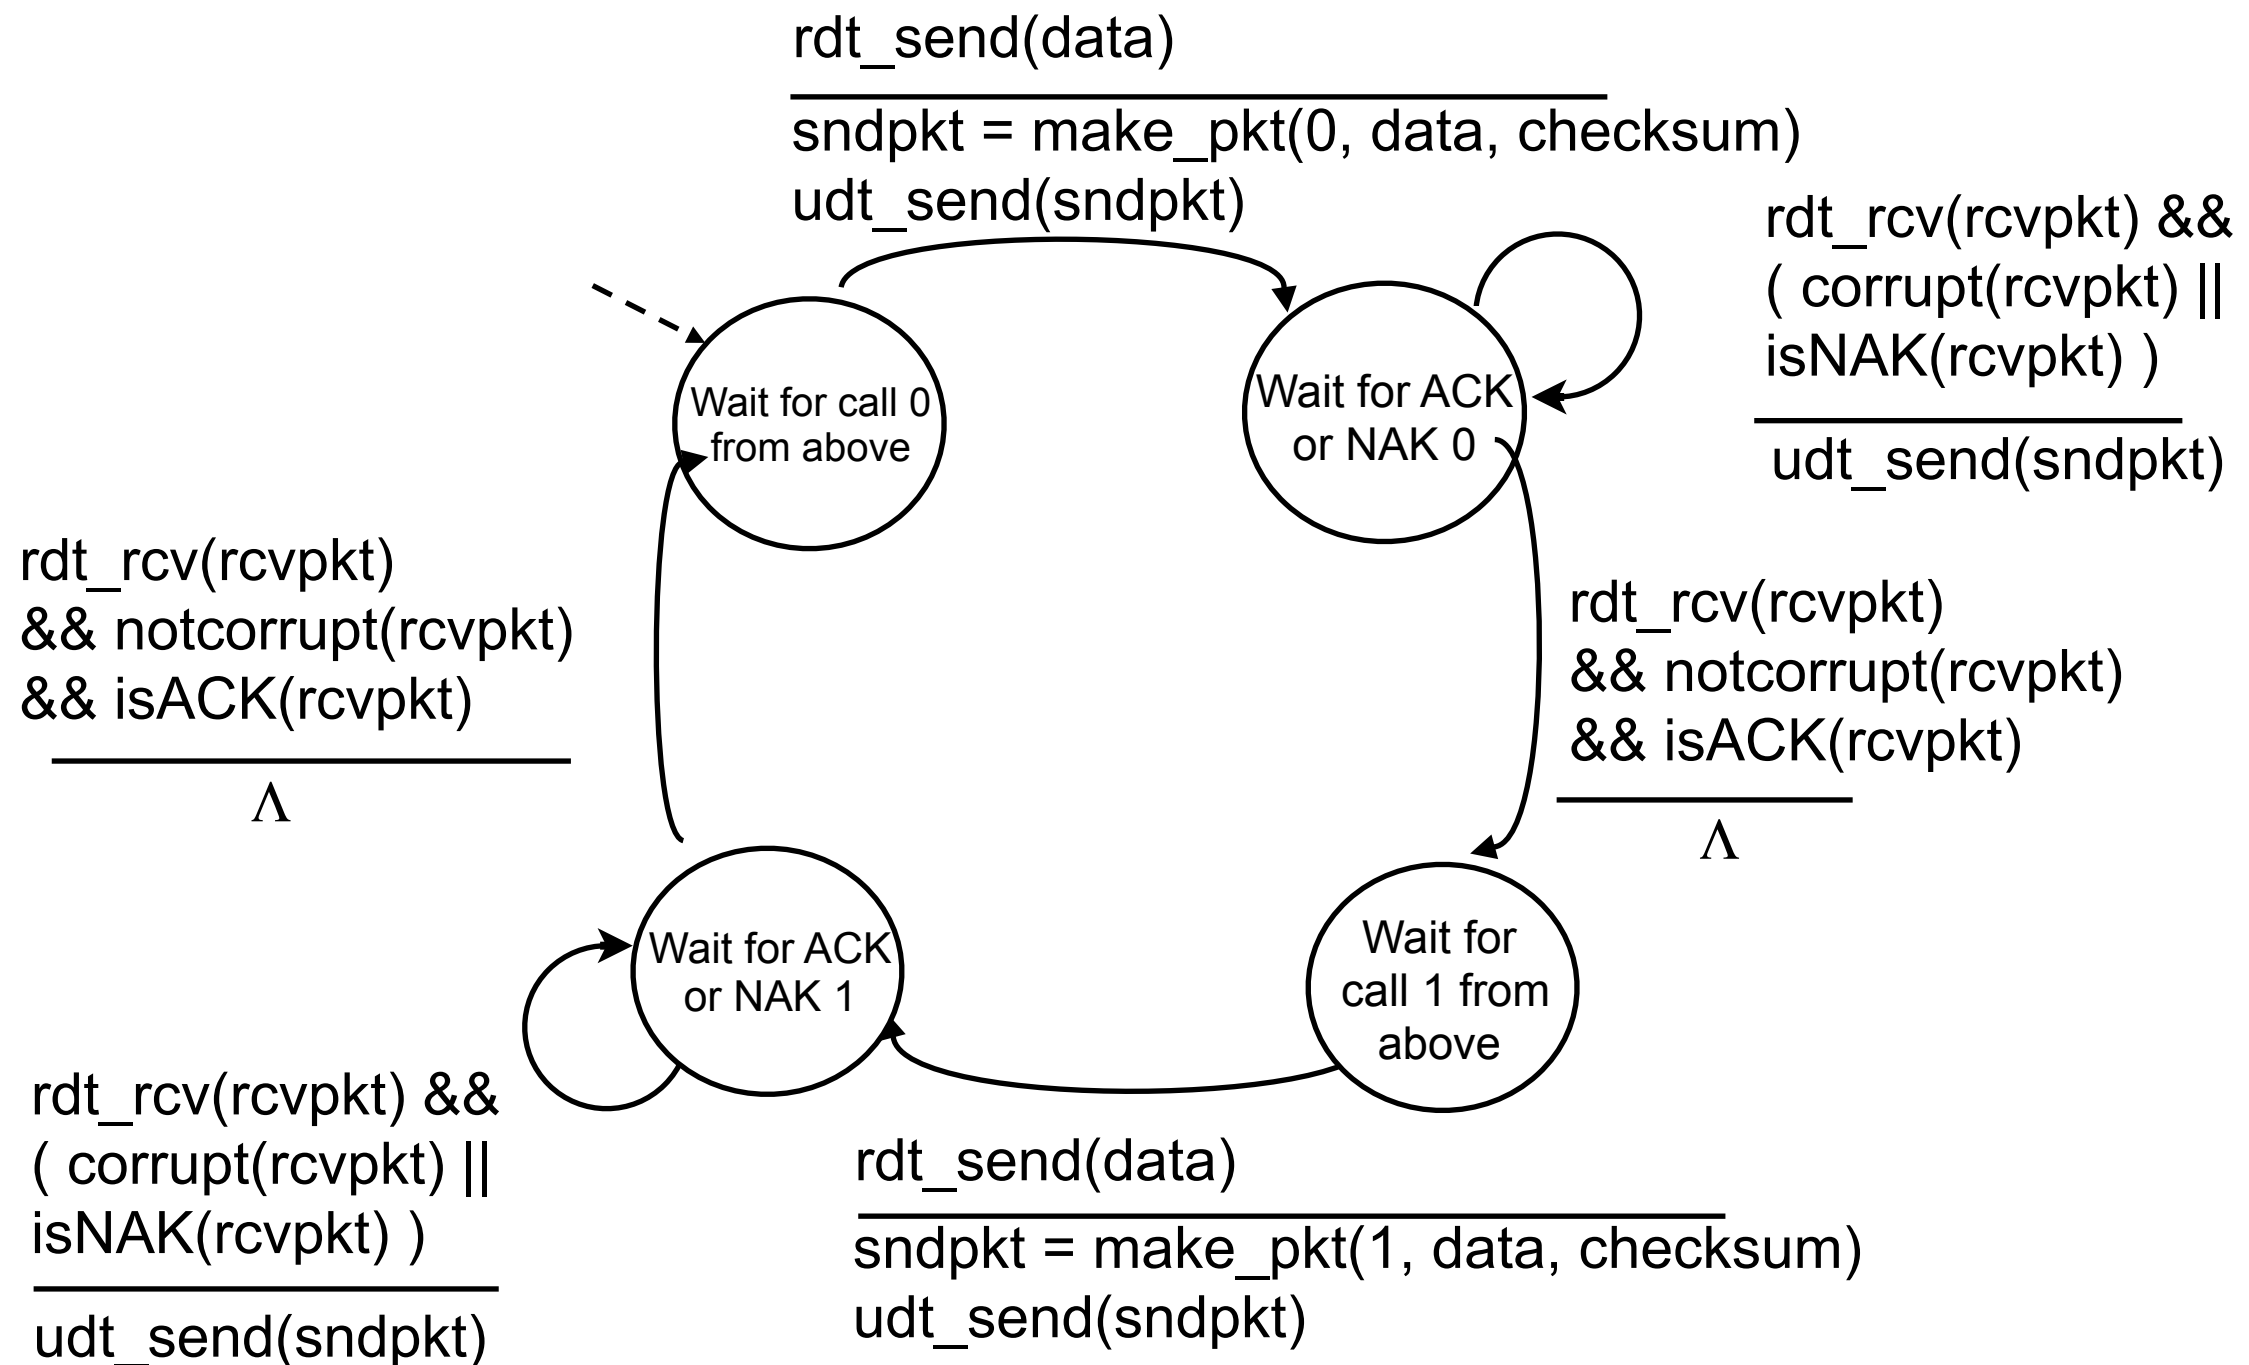
- sender retransmits current pkt if ACK/NAK garbled
- sender adds **sequence number** to each pkt
- receiver discards (doesn't deliver up) duplicate pkt



### stop and wait

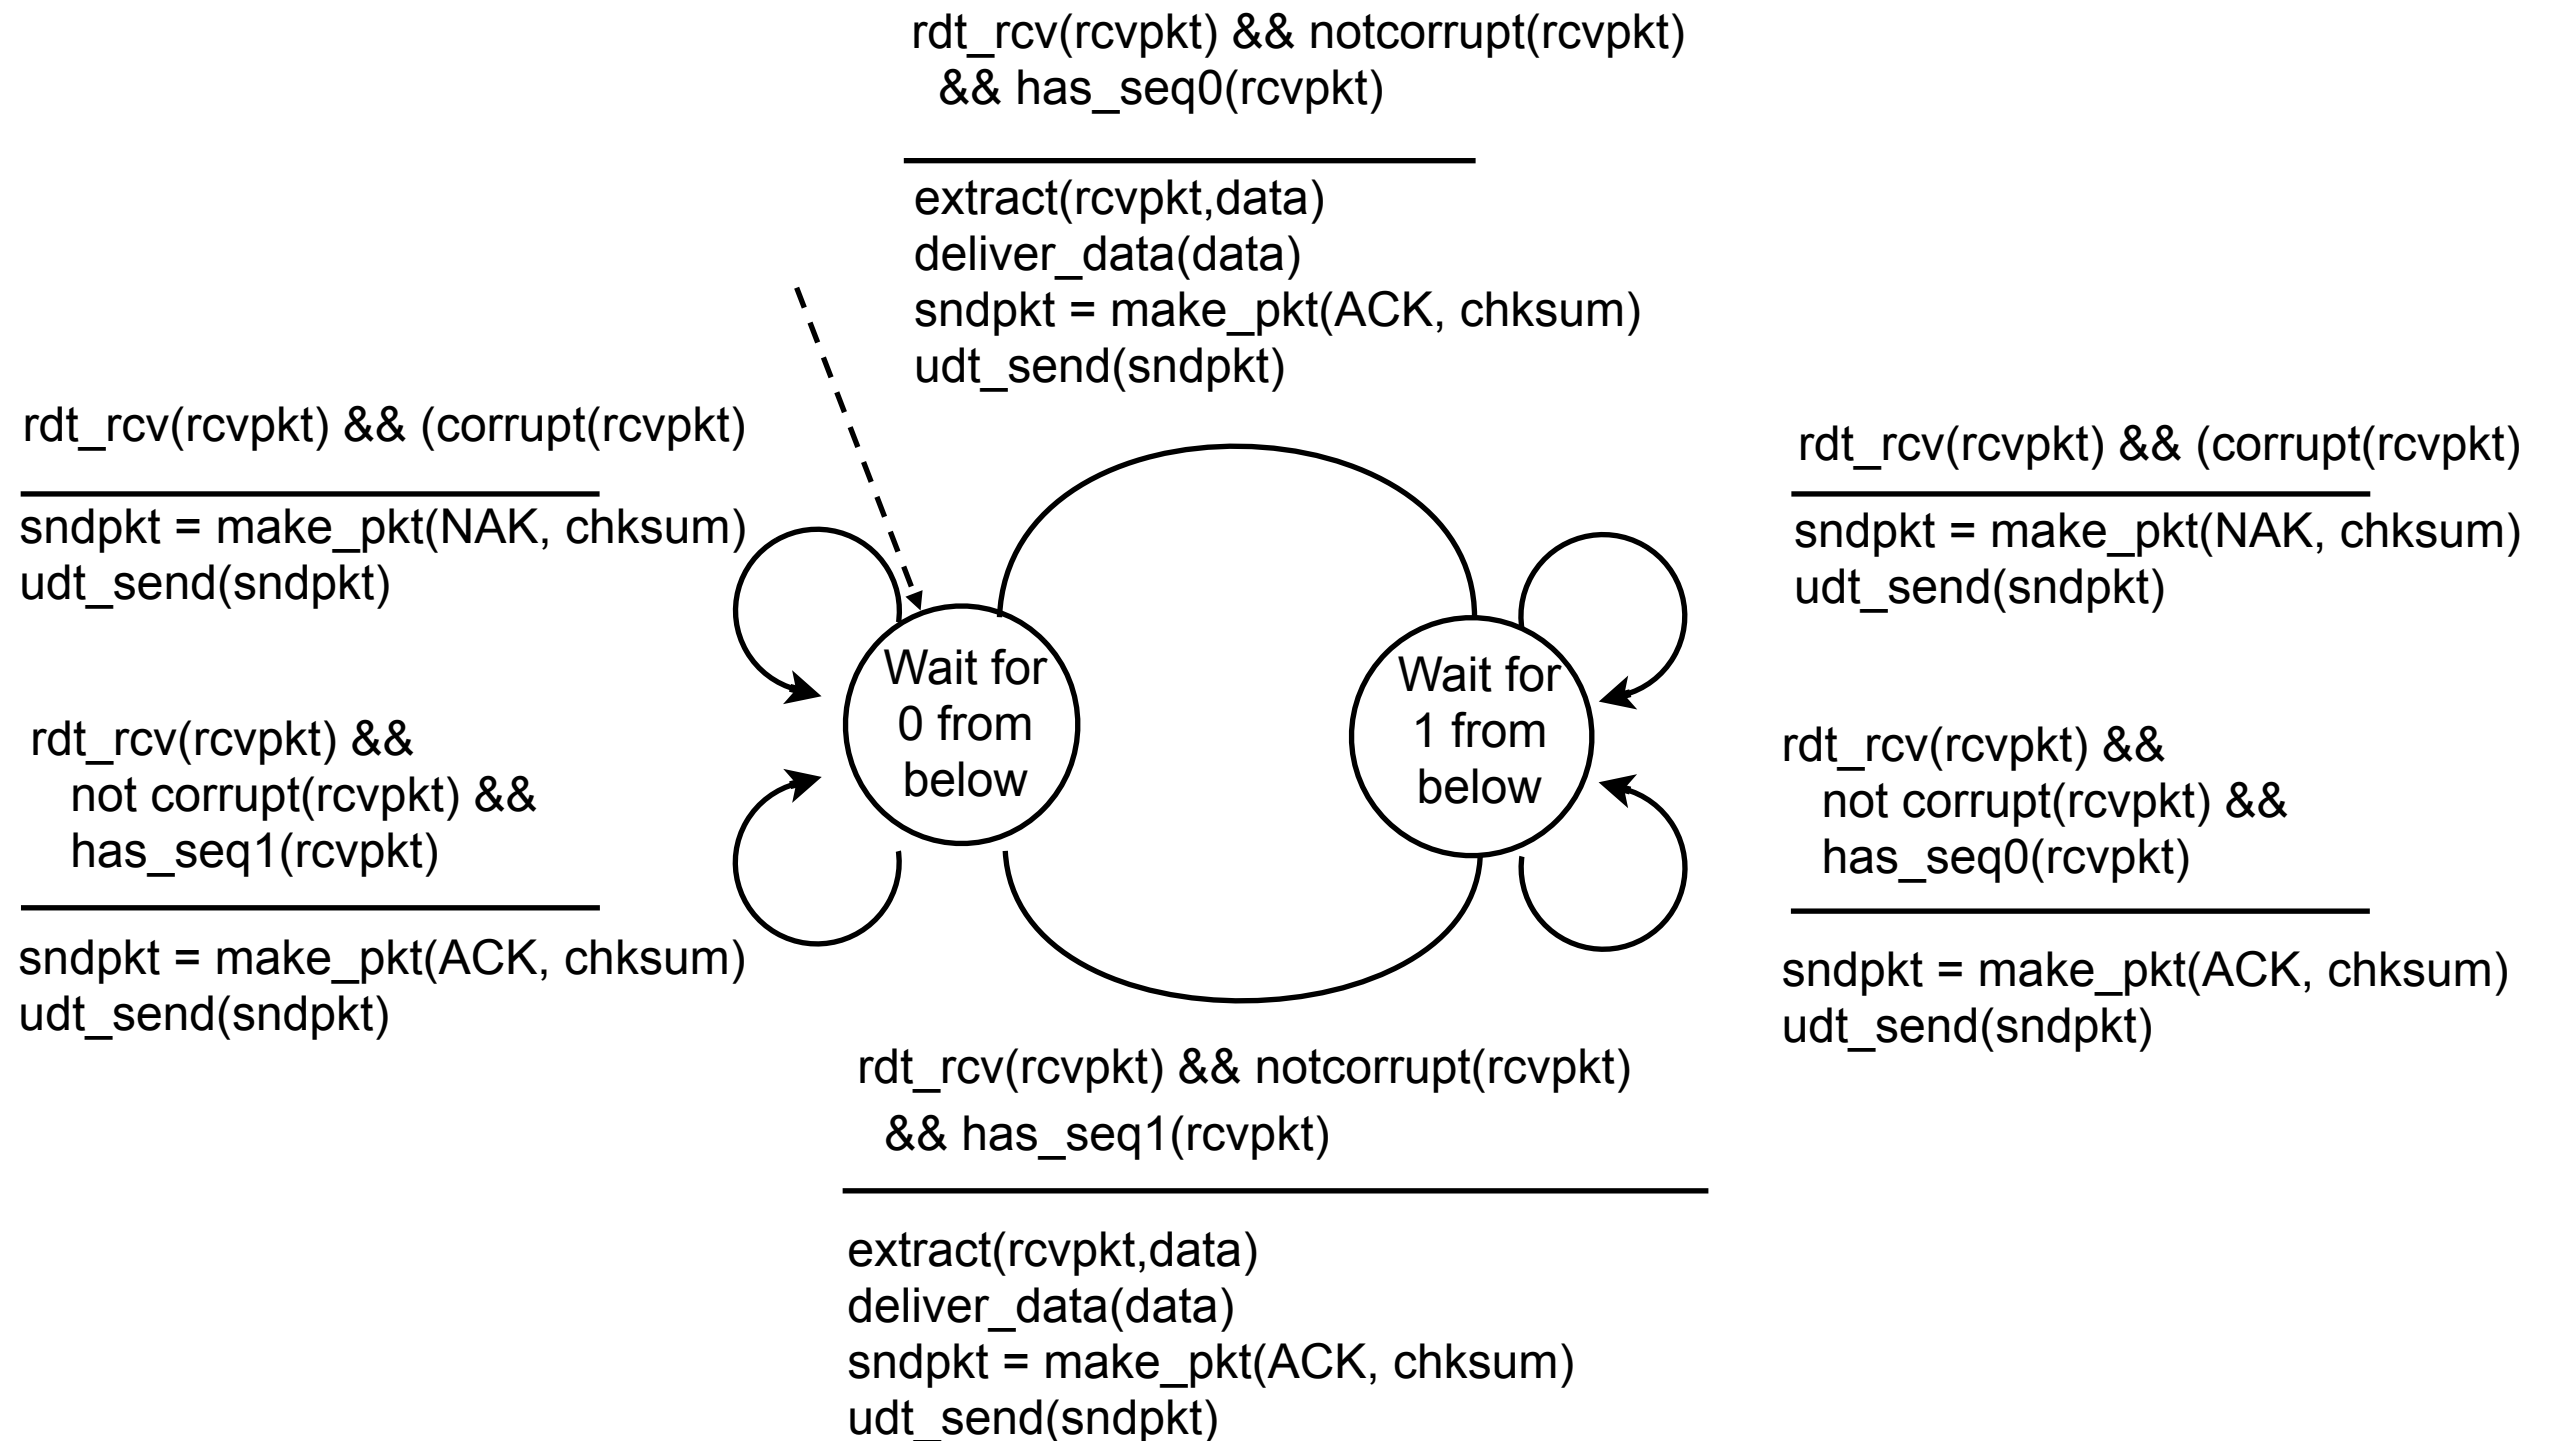
Sender sends one packet,  
then waits for receiver  
response

# rdt2.1: sender, handles garbled ACK/NAKs





# rdt2.1: receiver, handles garbled ACK/NAKs



# rdt2.1: discussion

## Sender:

- seq # added to pkt
- two seq. #'s (0, 1) will suffice.  
Why?
- must check if received ACK/NAK corrupted
- twice as many states
  - state must “remember” whether “current” pkt has 0 or 1 seq. #

## Receiver:

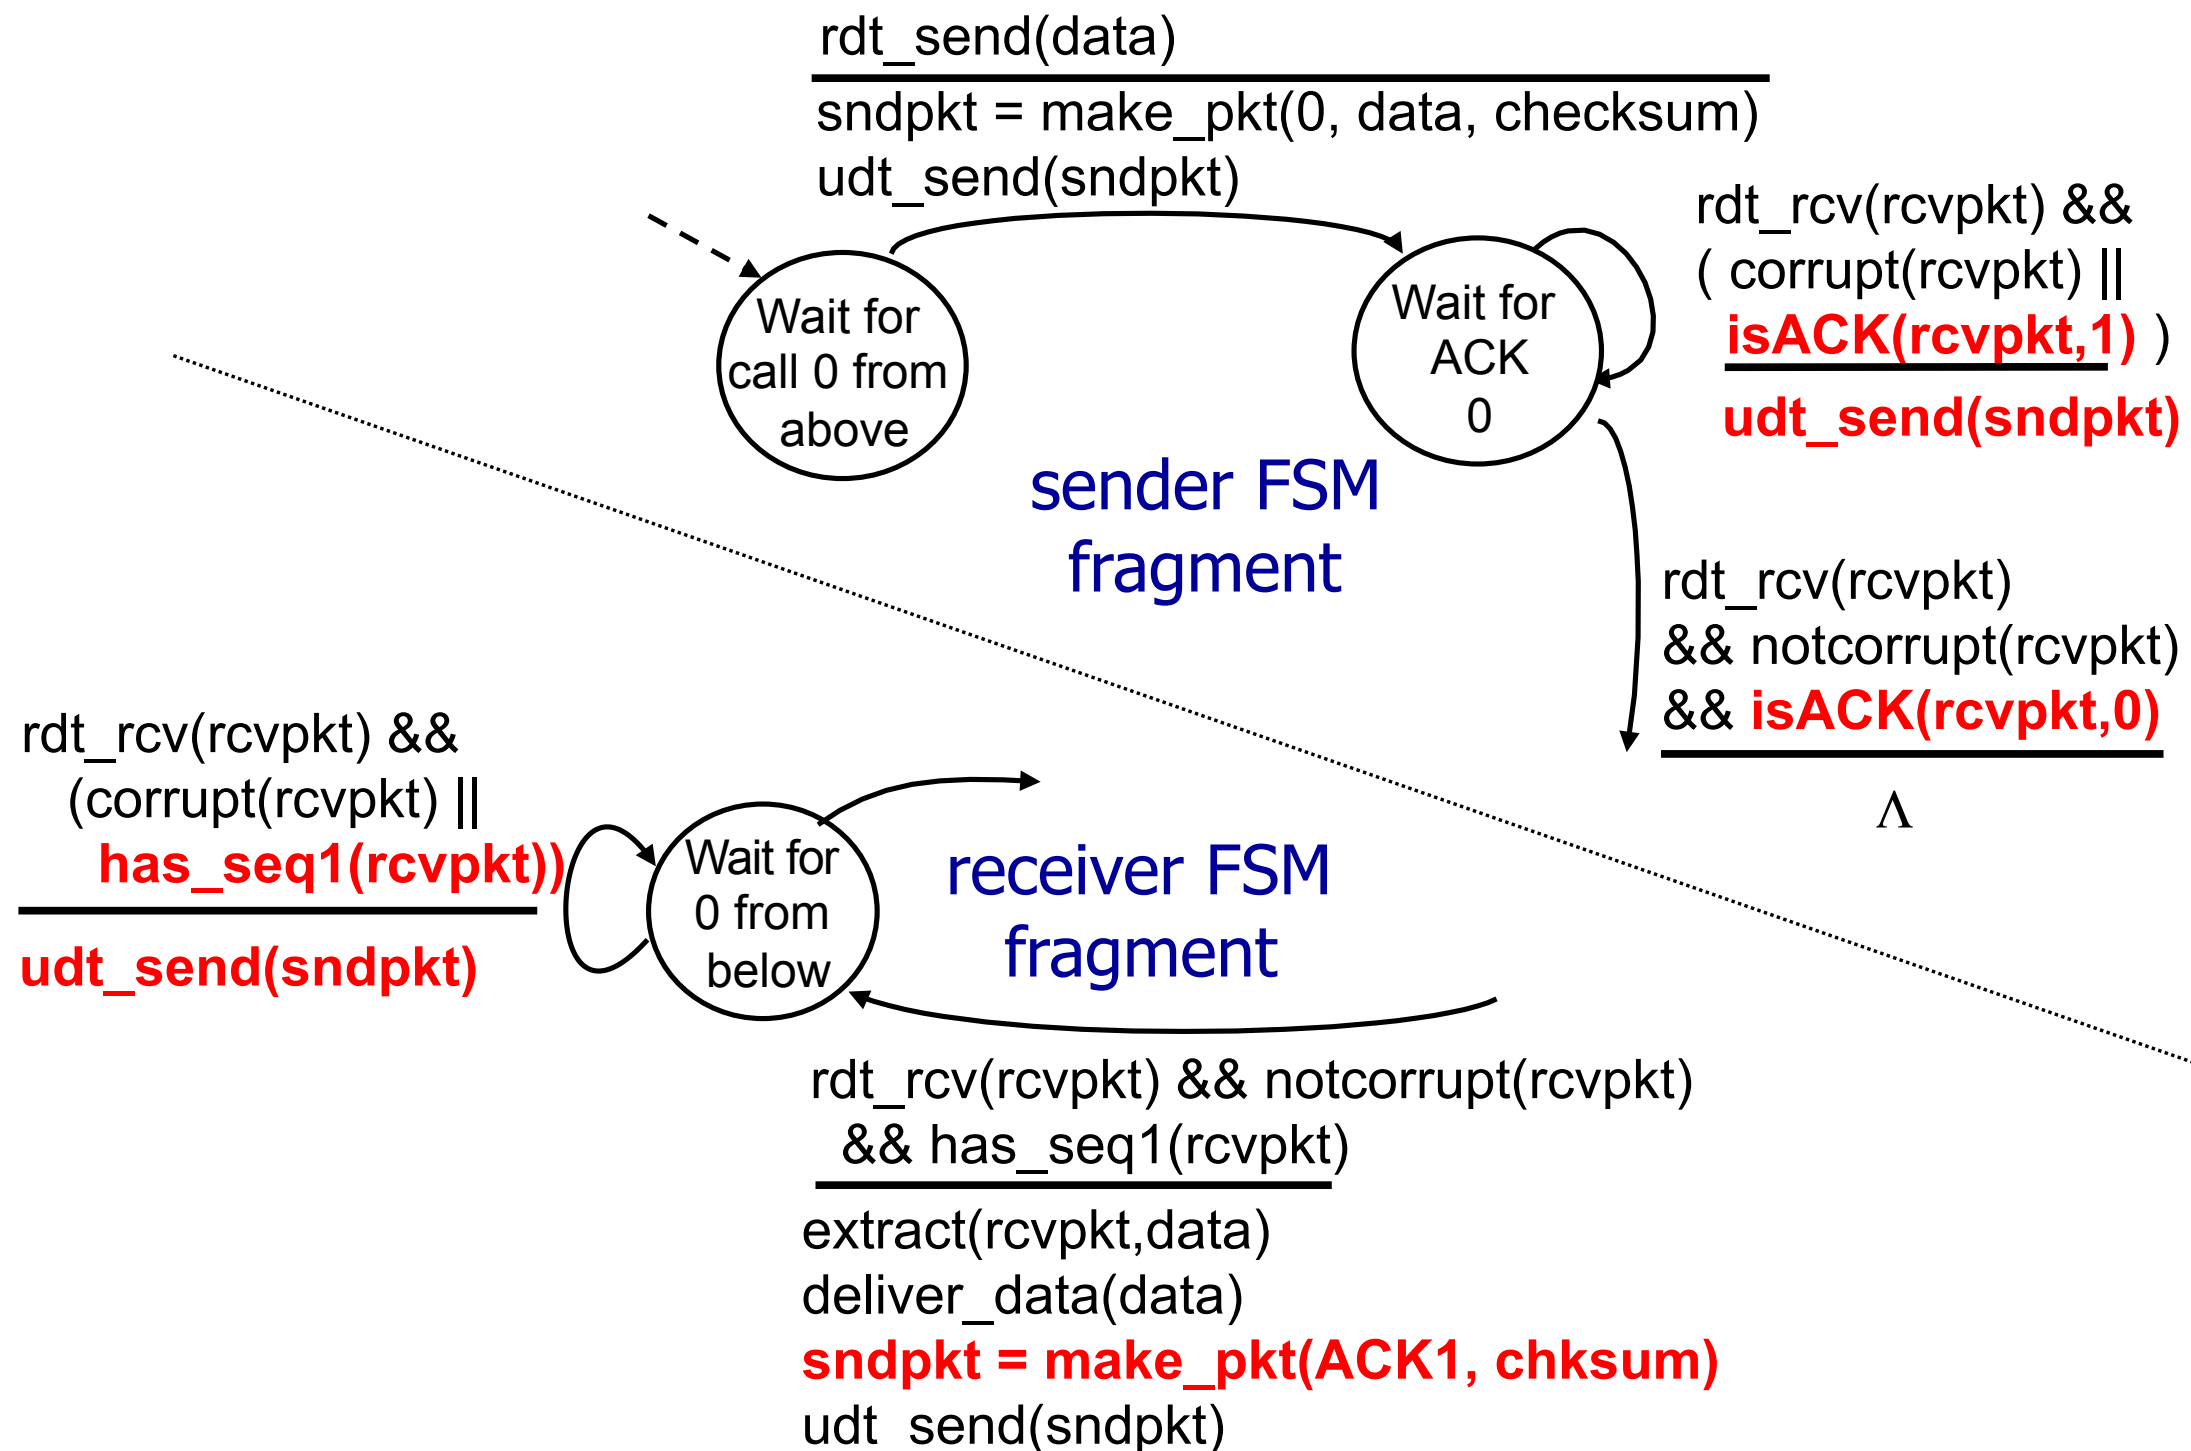
- must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can not know if its last ACK/NAK received OK at sender

# rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must explicitly include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*



# rdt2.2: sender, receiver fragments



# rdt3.0: channels with errors *and* loss

New assumption: underlying channel can also lose packets (data or ACKs)

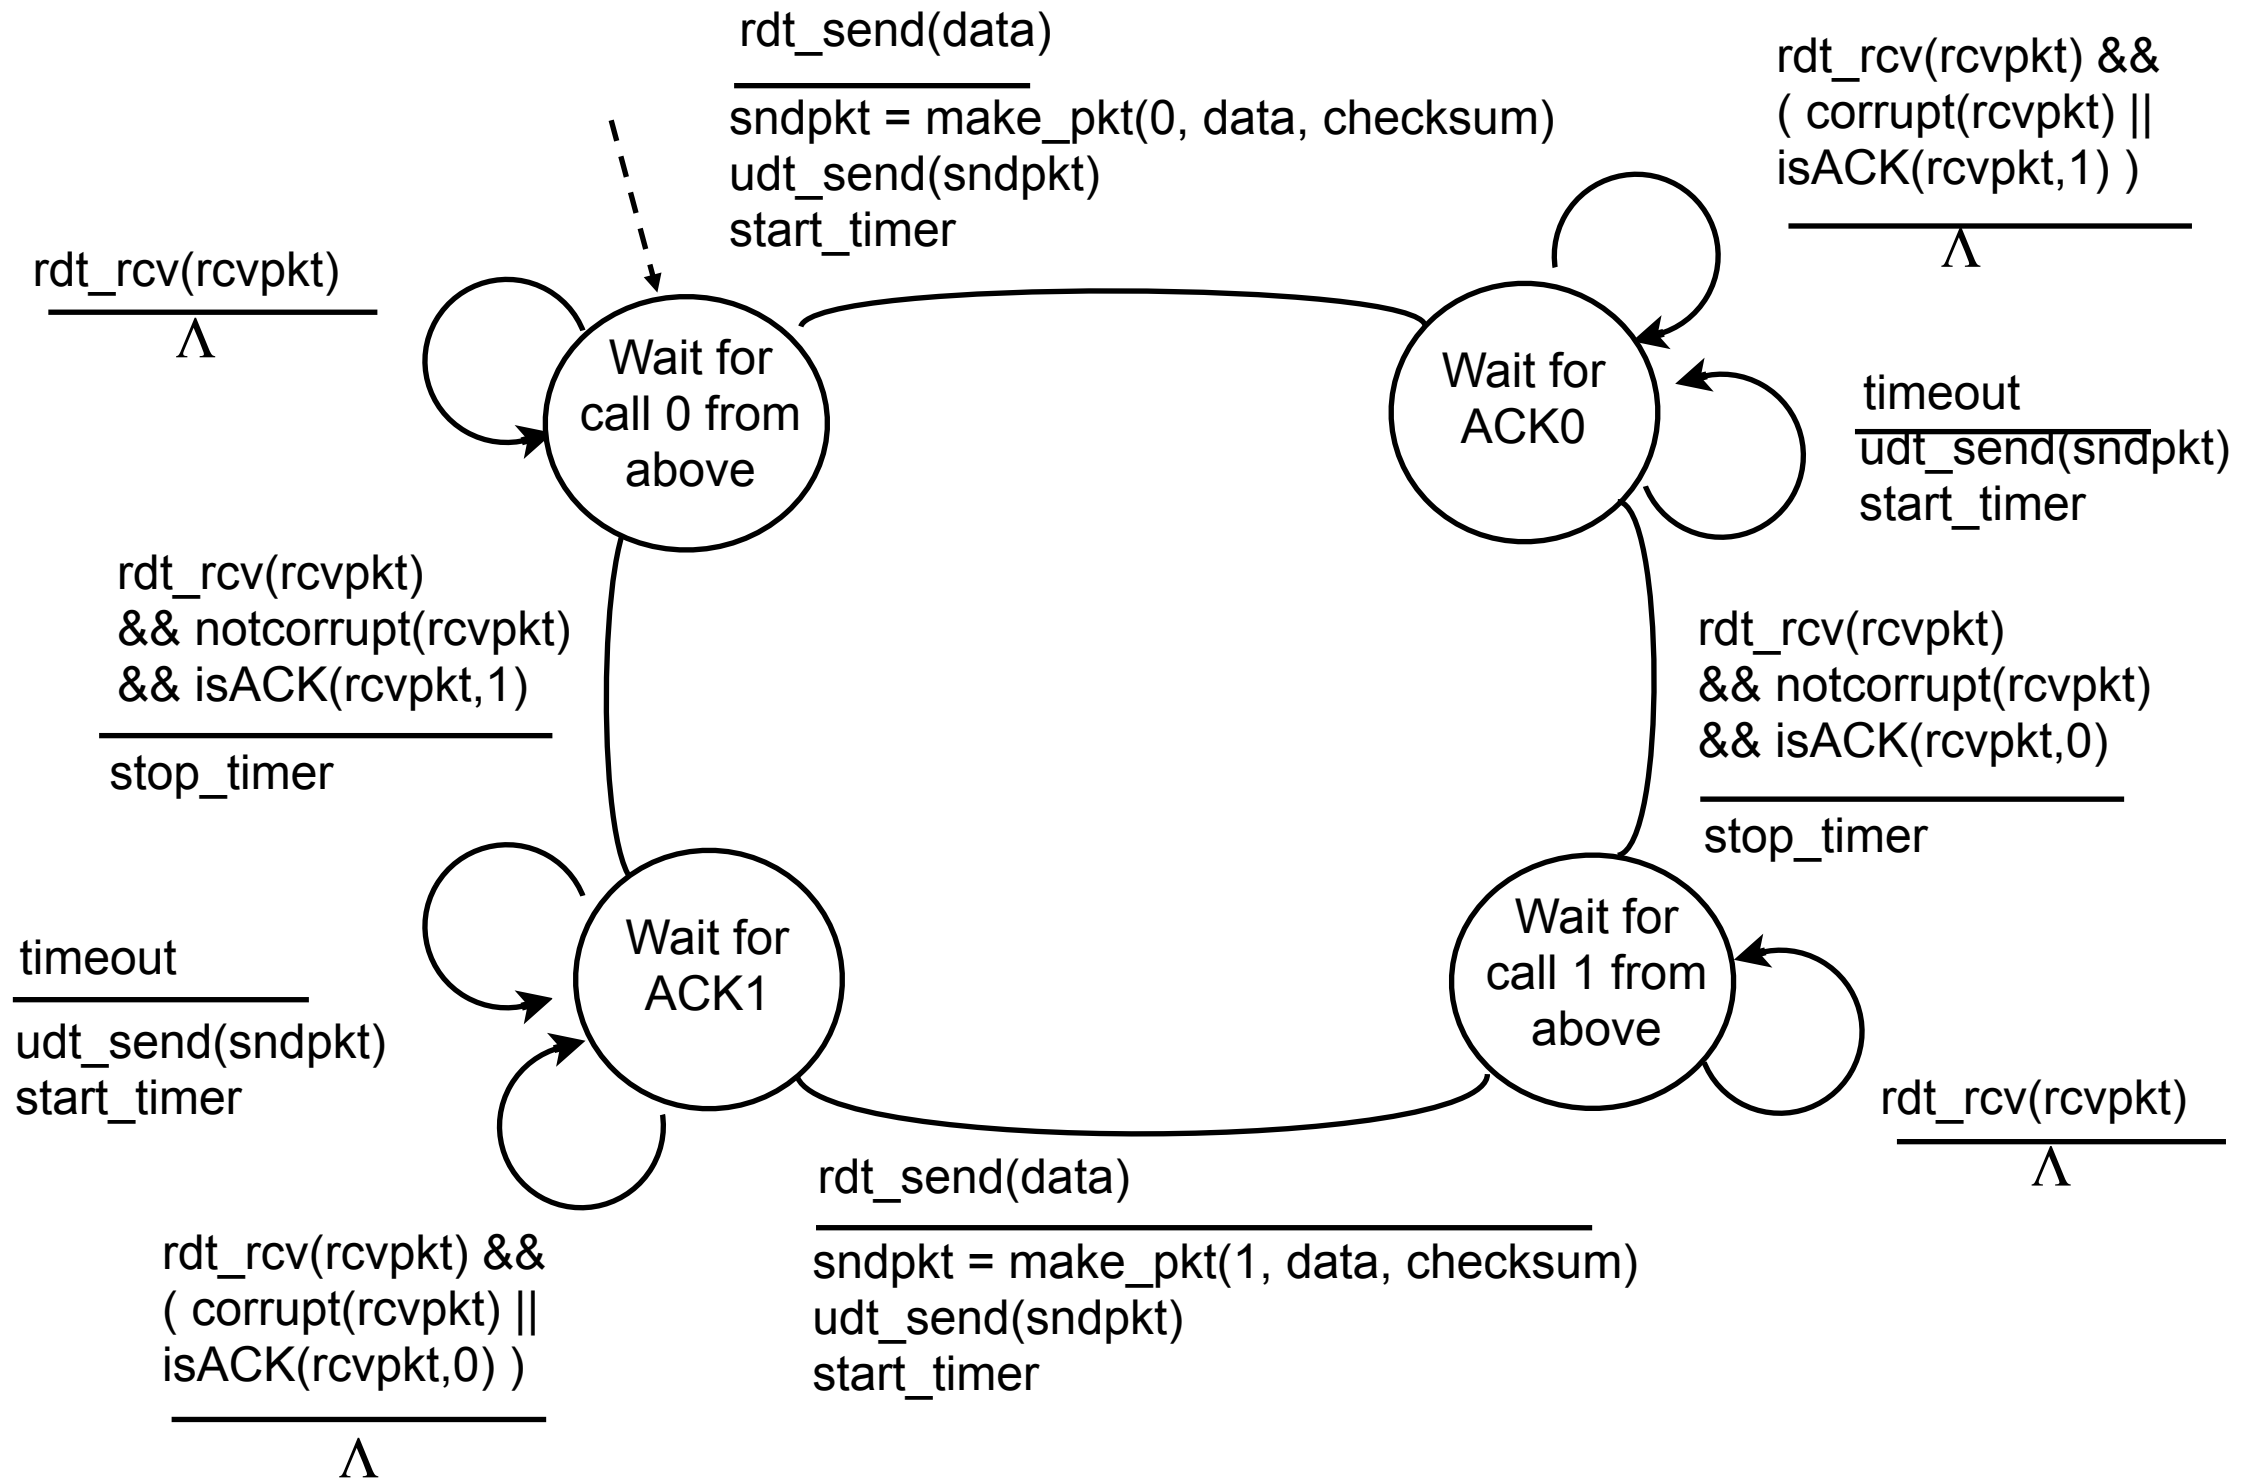
- ▶ checksum, seq. #, ACKs, retransmissions will be of help, but not enough



Approach: sender waits “reasonable” amount of time for ACK

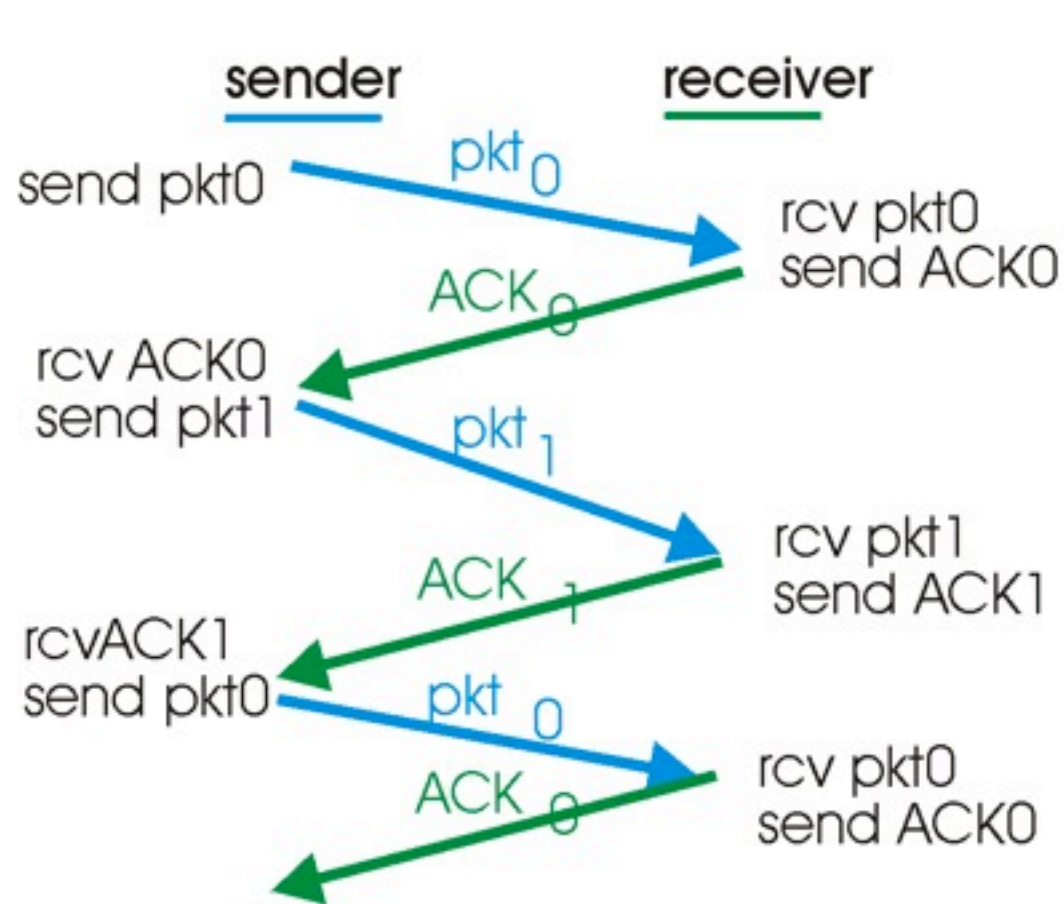
- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - ▶ retransmission will be duplicate, but use of seq. #'s already handles this
  - ▶ receiver must specify seq # of pkt being ACKed
- requires countdown timer

# rdt3.0 sender

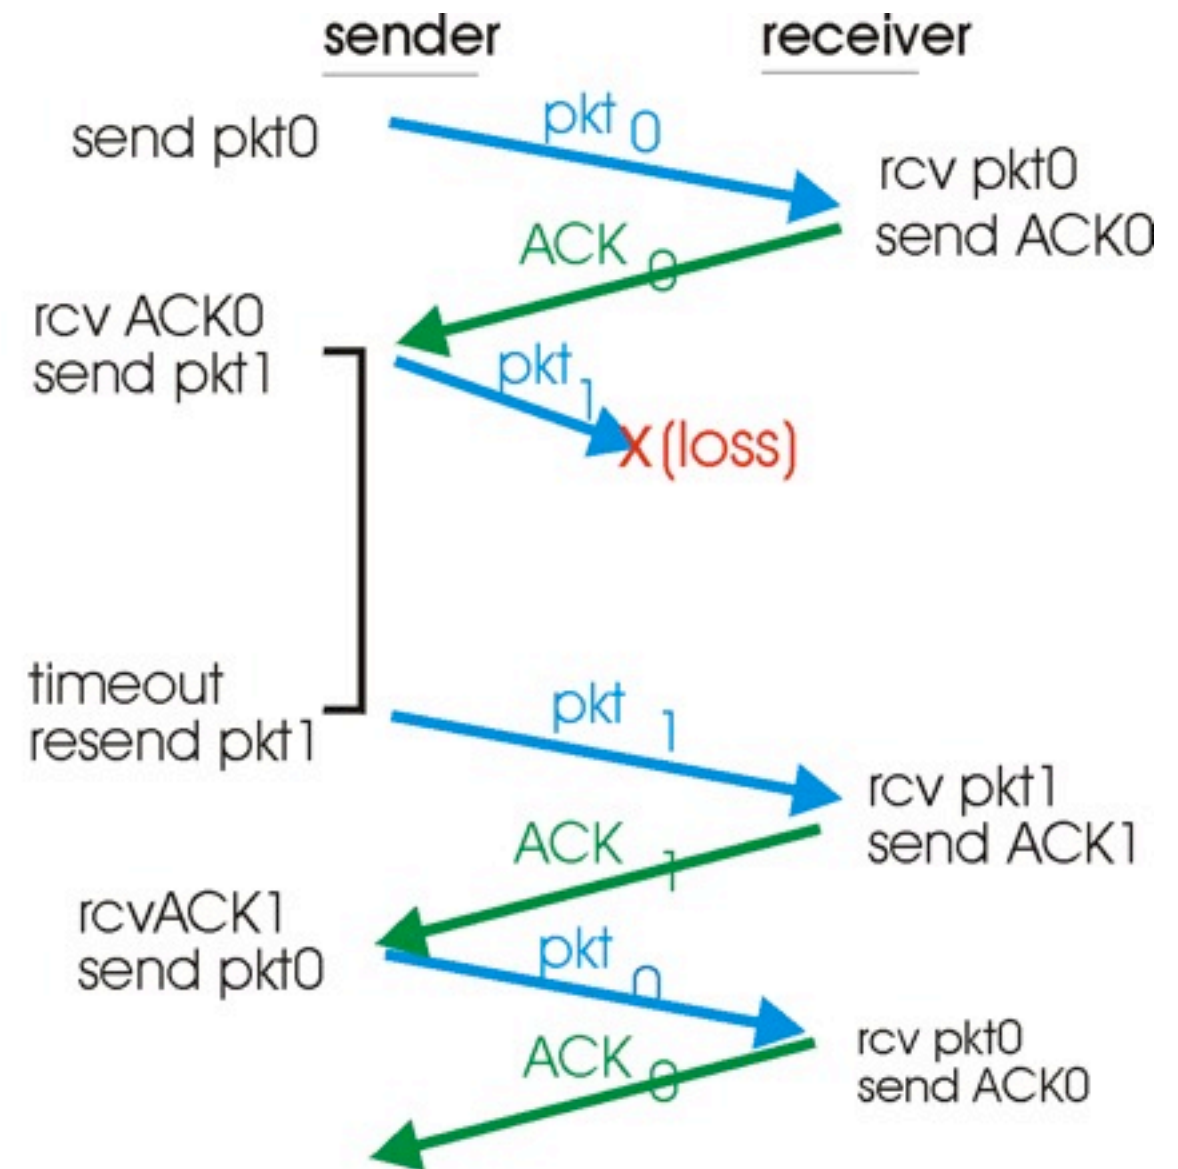




# rdt3.0 in action



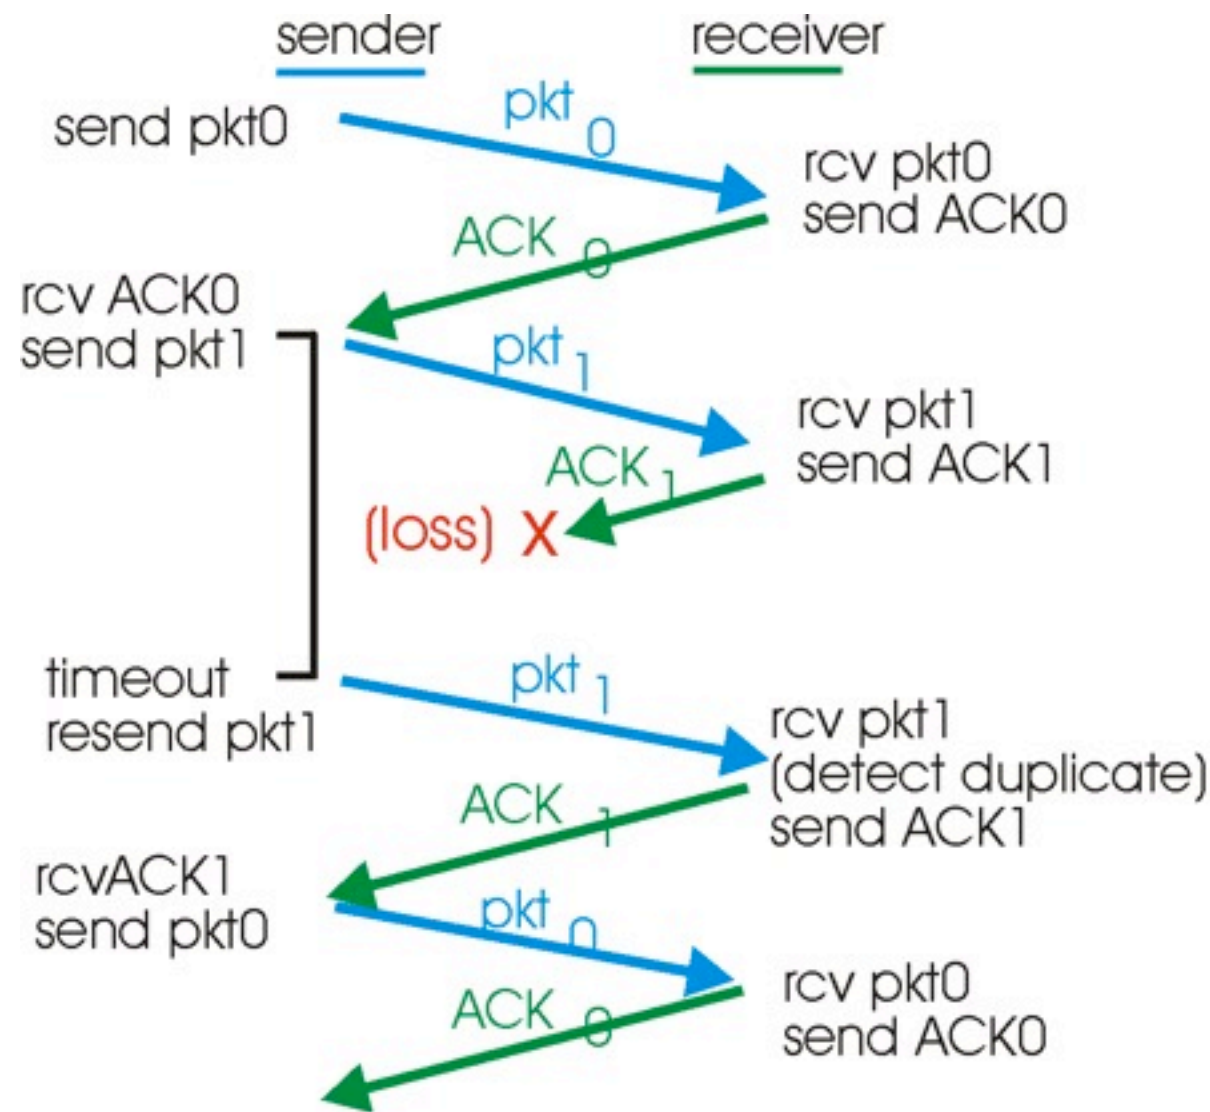
(a) operation with no loss



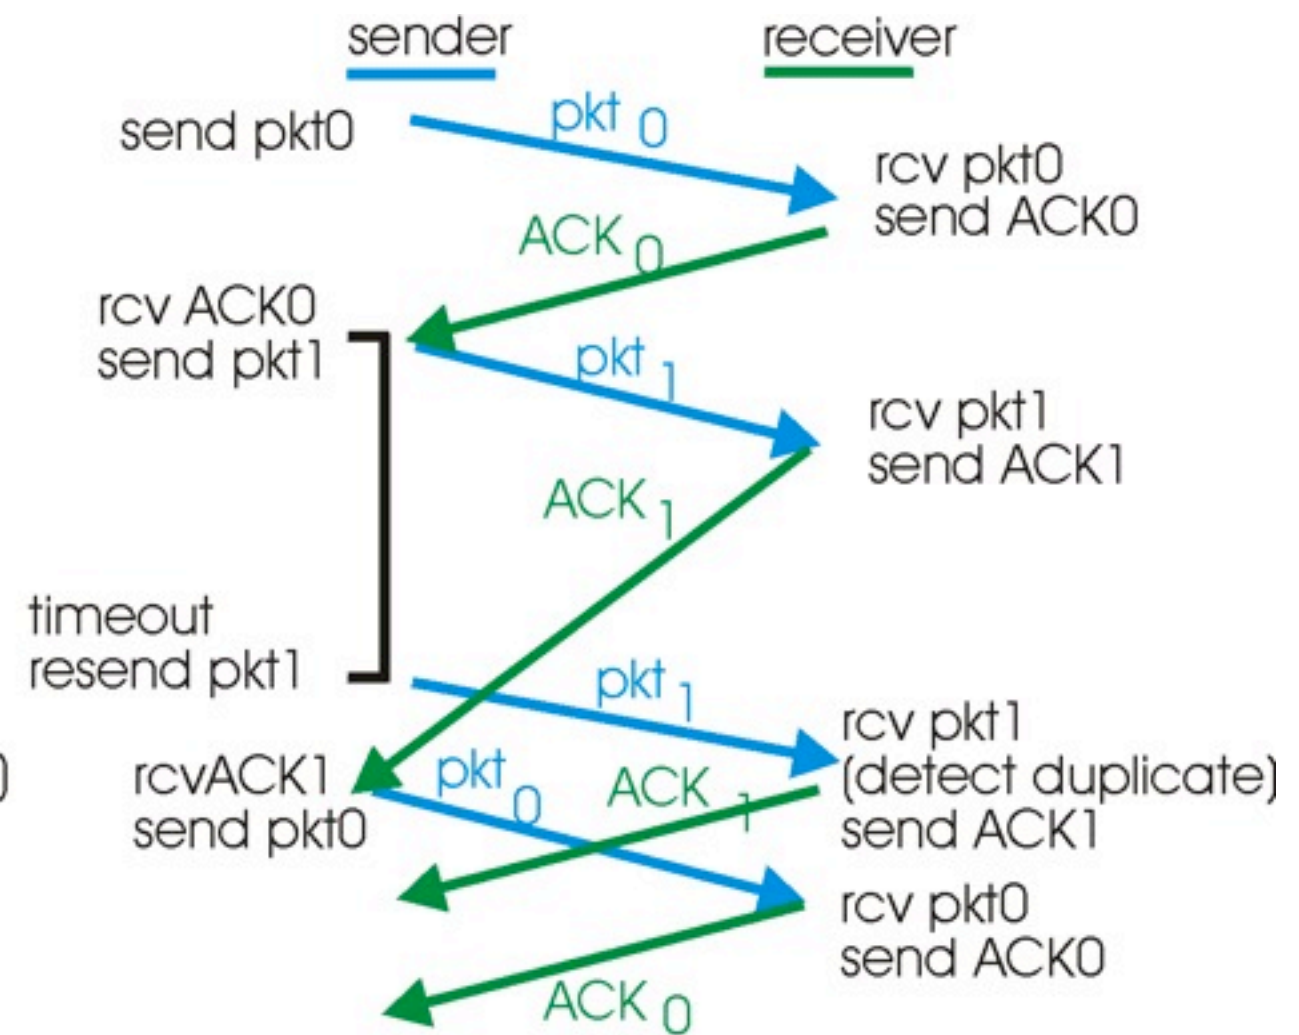
(b) lost packet



# rdt3.0 in action



(c) lost ACK



(d) premature timeout

# Performance of rdt3.0

- rdt3.0 works, but performance stinks
- example: 1 Gbps link, 15 ms e-e prop. delay, 1 KB packet:


$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^9 \text{ b/sec}} = 8 \text{ microsec}$$

- ▶  $U_{\text{sender}}$ : **utilization** – fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- ▶ 1 KB pkt every 30 msec -> 33kB/sec throughput over 1 Gbps link
- ▶ network protocol limits use of physical resources!

# rdt3.0: stop-and-wait operation




$$U_{sender} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

# Pipelined protocols

**Pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- ▶ range of sequence numbers must be increased
- ▶ buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- Two generic forms of pipelined protocols:  
**go-Back-N, selective repeat**

# Pipelining: increased utilization



**Increase utilization  
by a factor of 3!**

$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.000081$$

# Pipelining Protocols

- Go-back-N: big picture:
  - Sender can have up to  $N$  un-ACKed packets in pipeline
  - Rcvr only sends cumulative acks
    - Doesn't ack packet if there's a gap
  - Sender has timer for oldest unacked packet
    - If timer expires, retransmit all unacked packets



# Selective Repeat: Big Picture

- Sender can have up to  $N$  unacked packets in pipeline
- Rcvr acks individual packets
- Sender maintains timer for each unacked packet
  - When timer expires, retransmit only individual unacked packet.






# Go-Back-N

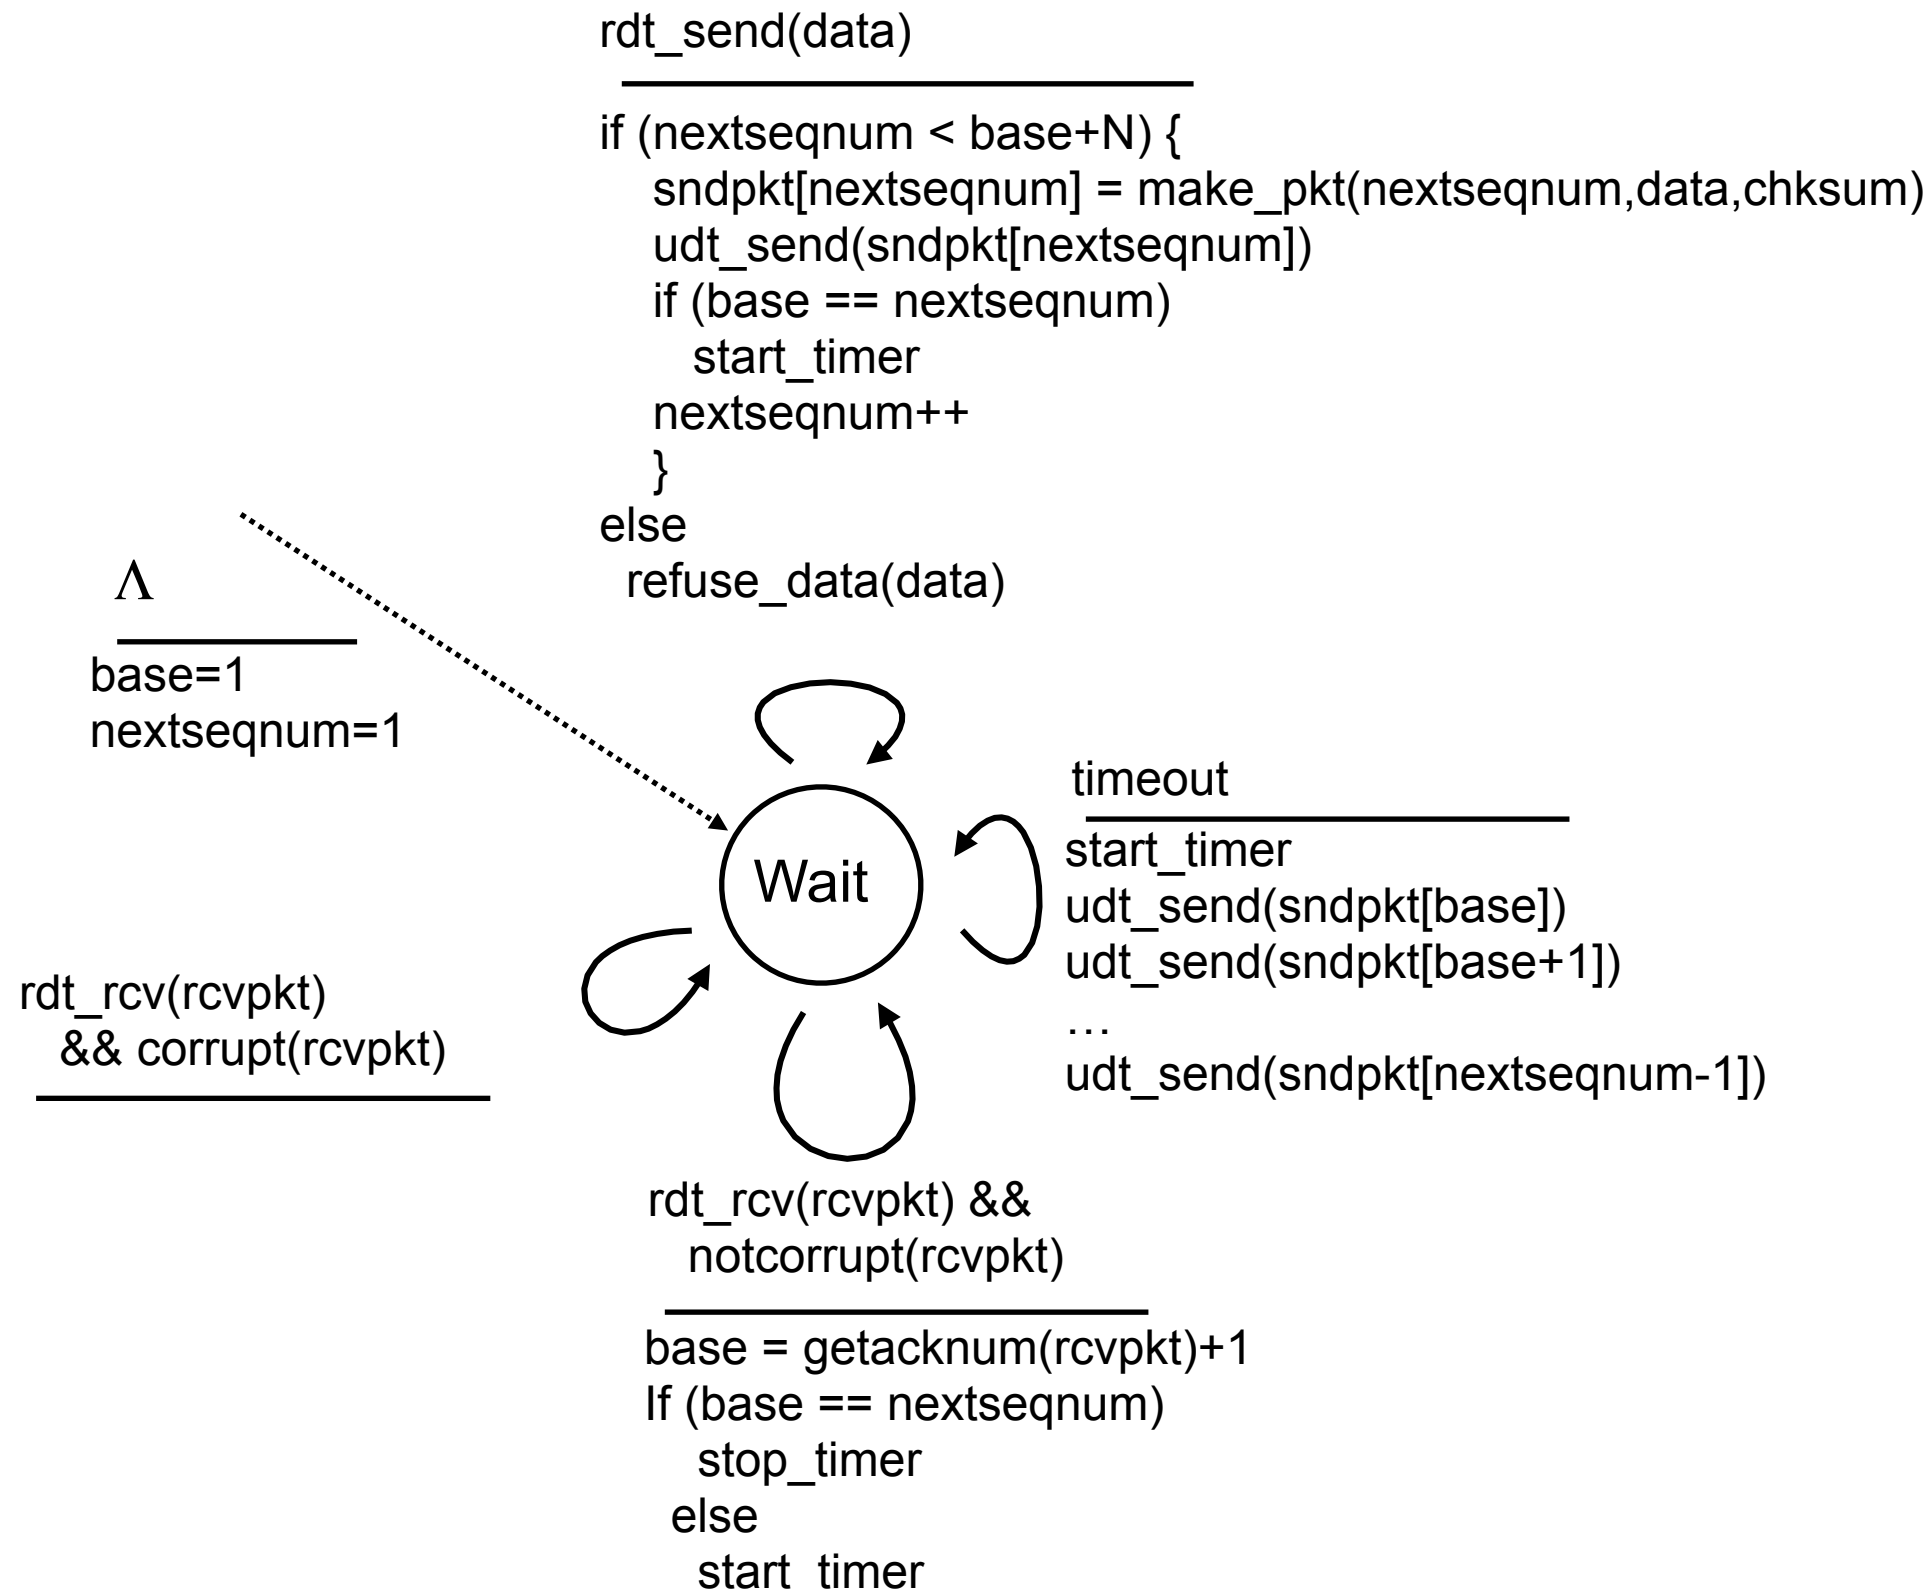
## Sender:

- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed

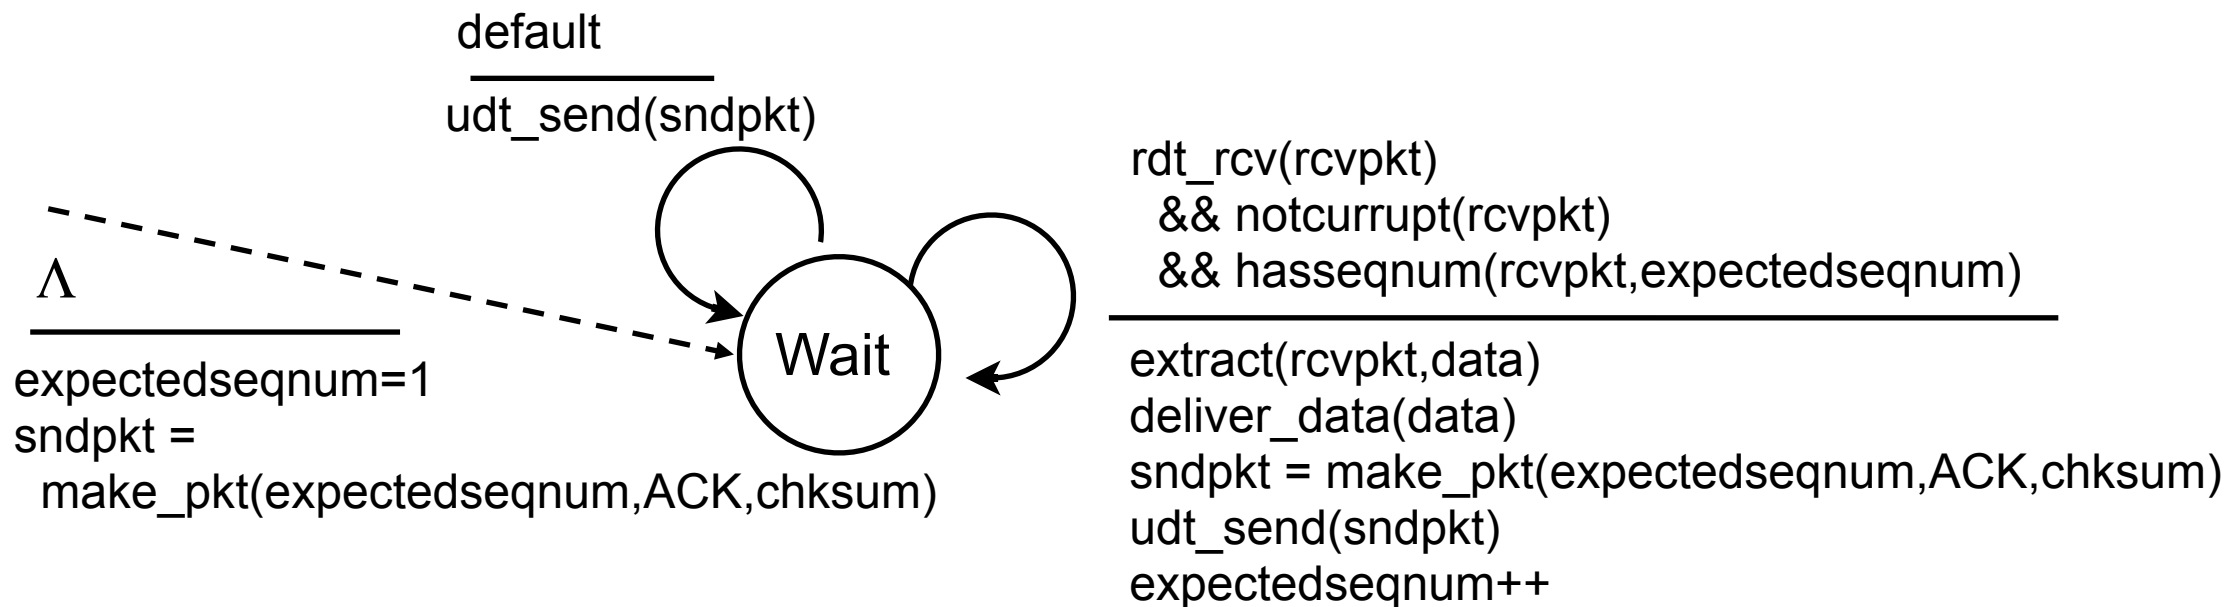


- ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”
  - ▶ may receive duplicate ACKs (see receiver)
- timer for each in-flight pkt
- timeout(n): retransmit pkt n and all higher seq # pkts in window

# GBN: sender extended FSM



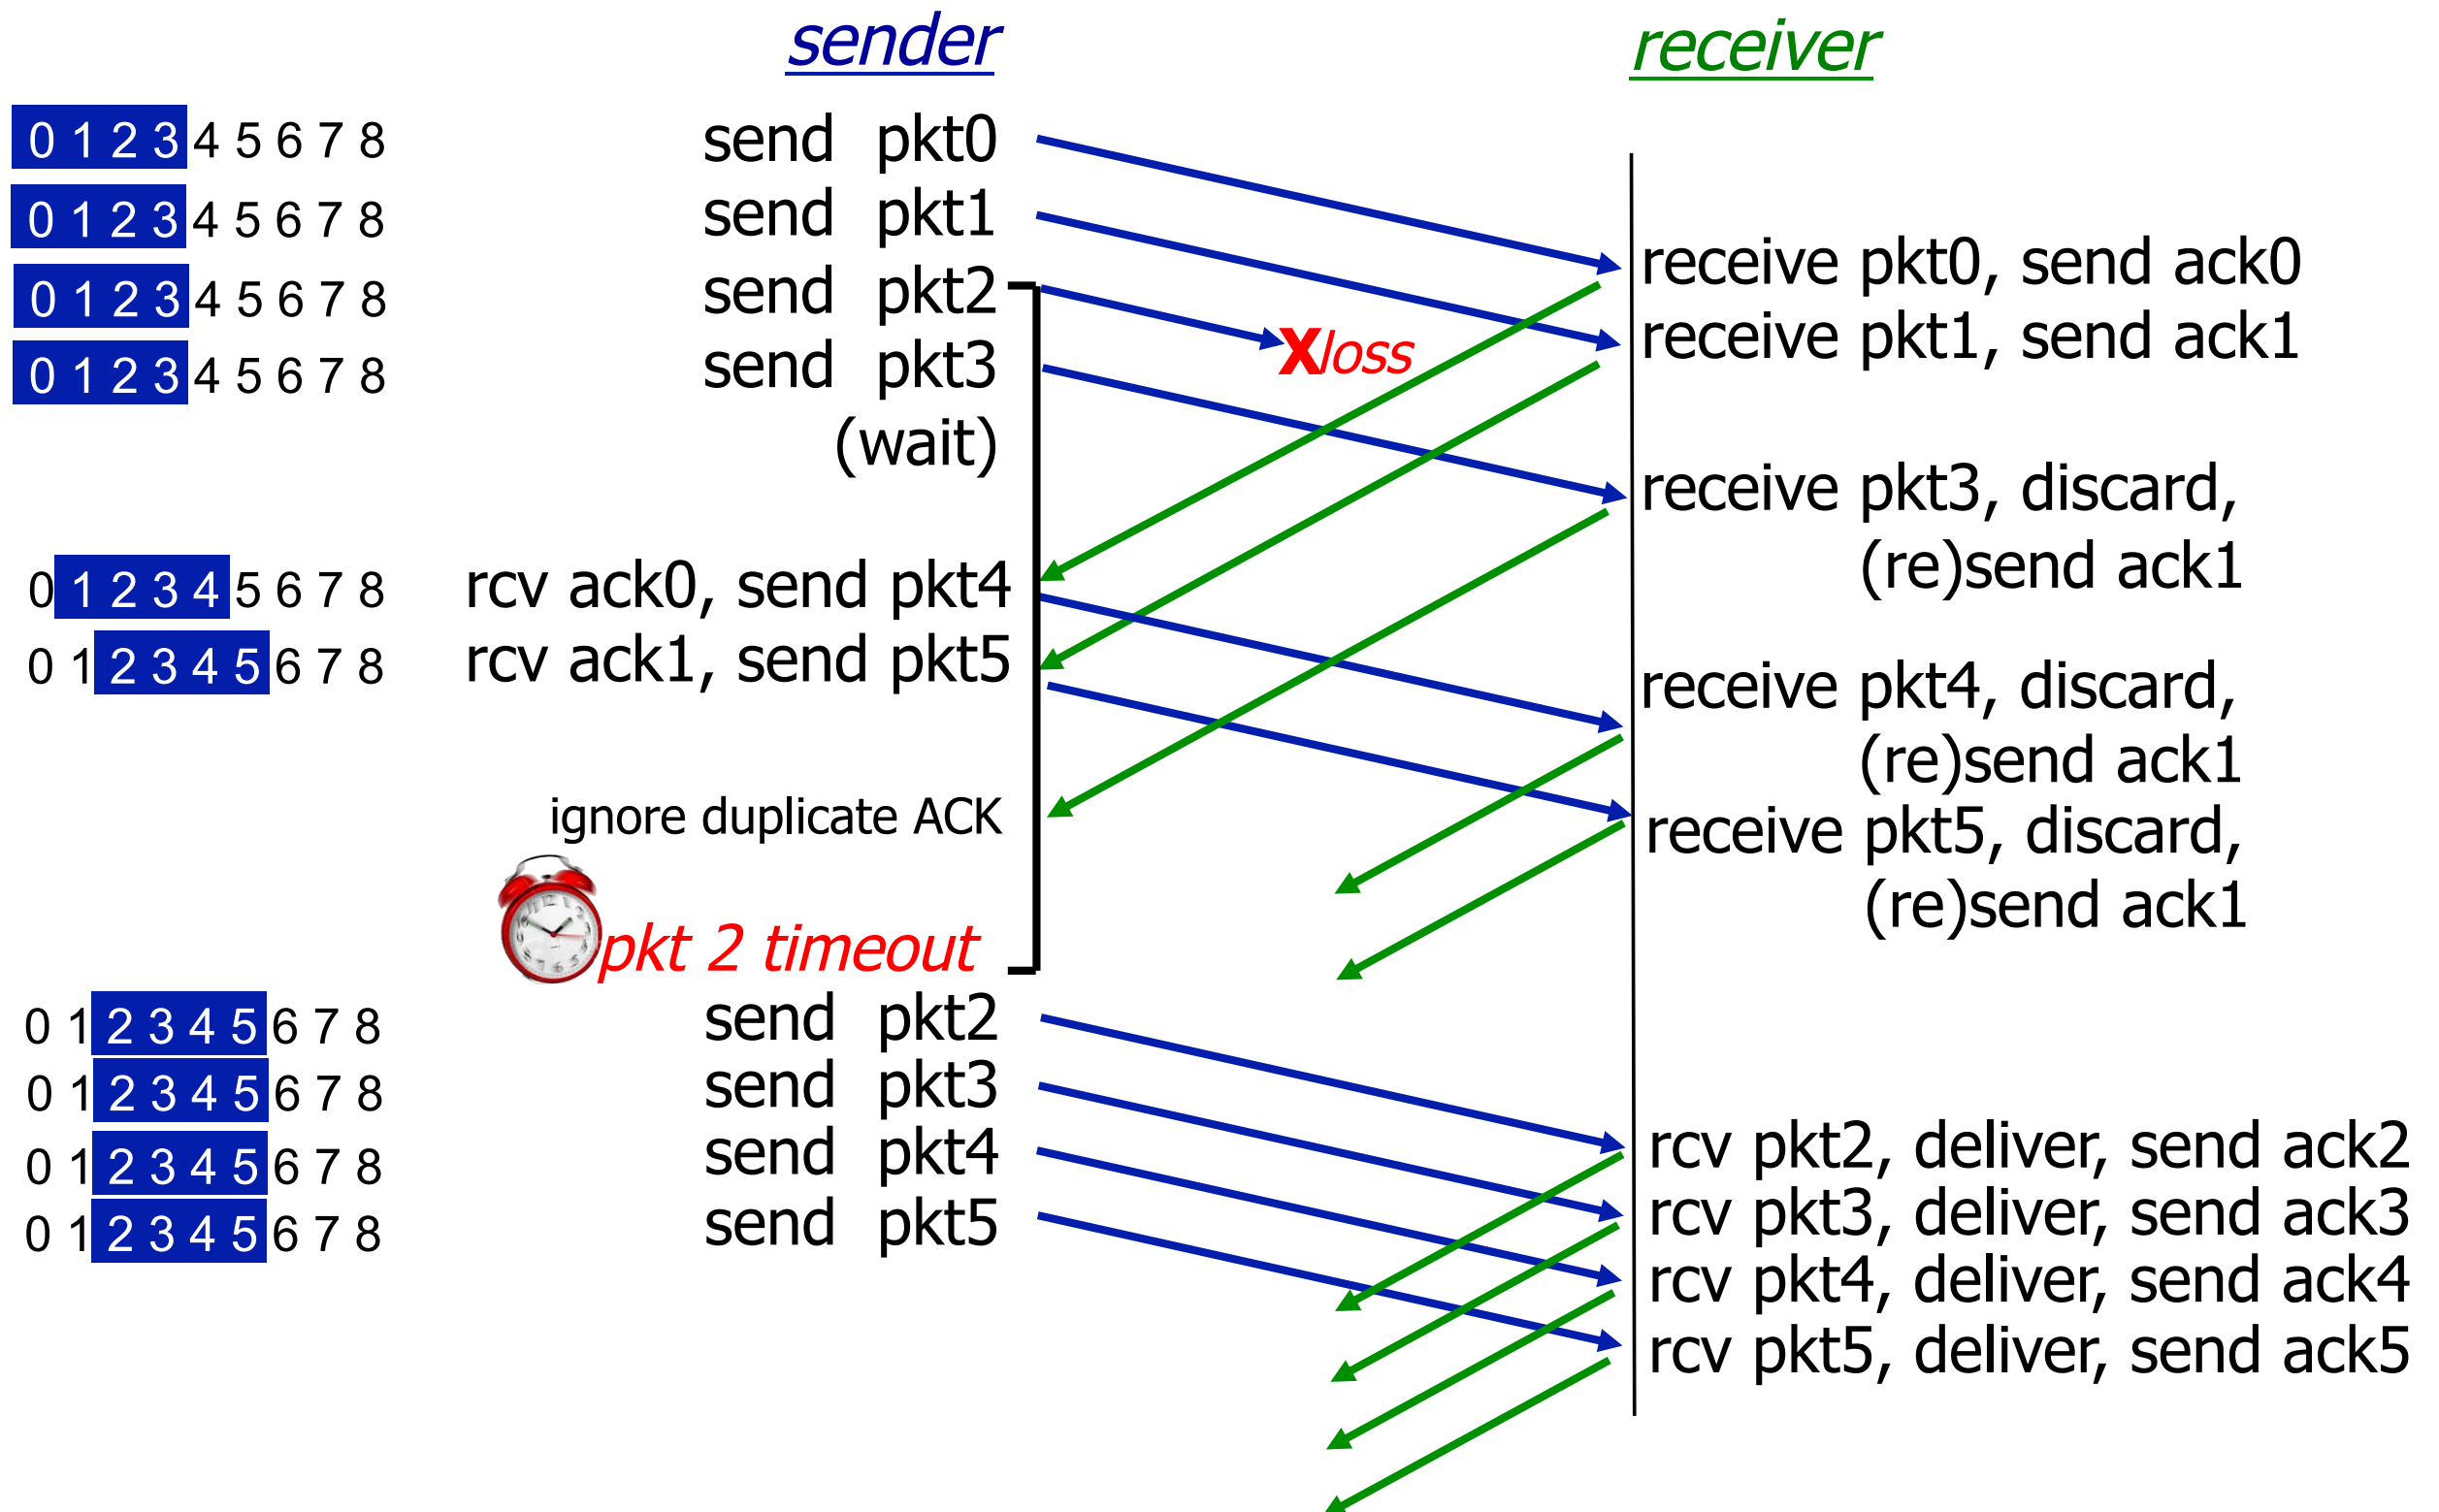
# GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest **in-order** seq #

- ▶ may generate duplicate ACKs
- ▶ need only remember **expectedseqnum**
- out-of-order pkt:
  - ▶ discard (don't buffer) -> **no receiver buffering!**
  - ▶ Re-ACK pkt with highest in-order seq #


# GBN in action



# Selective Repeat

- receiver *individually* acknowledges all correctly received pkts
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt
- sender window
  - N consecutive seq #'s
  - again limits seq #'s of sent, unACKed pkts

# Selective repeat: sender, receiver windows



# Selective repeat

## sender

### data from above :

- if next available seq # in window, send pkt

### timeout(n):

- resend pkt n, restart timer

### ACK(n) in [sendbase, sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

## receiver

### pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

### pkt n in [rcvbase-N, rcvbase-1]

- ACK(n)

### otherwise:

- ignore



# Selective repeat in action

sender window (N=4)

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

sender

send pkt0  
send pkt1  
send pkt2  
send pkt3  
(wait)

rcv ack0, send pkt4  
rcv ack1, send pkt5

record ack3 arrived



*pkt 2 timeout*

send pkt2

record ack4 arrived

record ack4 arrived

receiver

receive pkt0, send ack0  
receive pkt1, send ack1

receive pkt3, buffer,  
send ack3

receive pkt4, buffer,  
send ack4

receive pkt5, buffer,  
send ack5

rcv pkt2; deliver pkt2,  
pkt3, pkt4, pkt5; send ack2


*Q: what happens when ack2 arrives?*

# Selective repeat: dilemma


## Example:

- seq #'s: 0, 1, 2, 3
- window size=3
- receiver sees no difference in two scenarios!
- incorrectly passes duplicate data as new in (a)

**Q:** what relationship between seq # size and window size?



*receiver can't see sender side.  
receiver behavior identical in both cases!  
something's (very) wrong!*



# Next Time

- That was a lot of material...
  - Take time to look over the notes. Understand the differences between each of these schemes!
- Next Time
  - TCP and Congestion Control (Sections 3.5 and 3.6)
- Check the webpage for Homework 2 and (soon) Project 2!

